The Fundamental Memory Ceiling of Noisy Quantum Reservoirs

Markus Baumann [1,*] Jonas Stein [1] and Claudia Linnhoff-Popien [1]

[1] *Quantum Applications and Research Laboratory (QAR-Lab), Chair of Mobile and Distributed Systems, Department of Computer Science, Ludwig-Maximilians-Universität München, Oettingenstraße 67, 80538 München, Germany*

(Dated: October 18, 2026)

We show that enlarging a noisy quantum reservoir does not raise its fundamental memory ceiling. With k input qubits reinitialized per time step, the longest memory any reservoir can attain under depolarizing noise is $\exp[\Theta(k)]$ time steps, whatever its size, dynamics, and input encoding. A universal upper bound and an explicit reservoir that attains the exponential scaling establish this optimal resource law. It extends to an entropy budget for general noise and applies to any noisy quantum dynamics refreshed by fresh qubits, including measurement-free error correction.

Introduction.—Reservoir computing processes time series with the transient dynamics of a fixed physical system, the reservoir: inputs are injected into it, and only a linear readout is trained [1–5]. Quantum reservoir computing (QRC) uses interacting qubits for this purpose [6–8] and has been realized on superconducting processors with up to 124 qubits [9–14]. For temporal tasks, memory is the decisive resource. A reservoir that forecasts a time series can use an input only as long as that input still influences its state, and on noisy hardware this influence fades at every time step. Can a noisy quantum reservoir, or any noisy quantum system refreshed by fresh qubits [Fig. 1], remember for arbitrarily long if it is only made large enough?

The answer is not obvious. A larger reservoir offers more degrees of freedom in which to distribute information about past inputs, potentially making it more robust to noise. Yet redundancy alone does not prevent information from gradually degrading. Preserving it over long times may also depend on how the system is continually refreshed: in QRC, the input qubits are discarded and reinitialized at every time step, providing a route for removing accumulated entropy [Fig. 2(b)] [6, 11, 12]. Noise brings entropy into every qubit at every time step, whereas only the reinitialized qubits carry entropy out.

Echo-state and fading-memory criteria guarantee that old inputs eventually lose their influence [1, 15–17], Markovian reservoirs forget exponentially fast [18], and dissipation and resets are recognized resources for QRC [9, 13, 19–21]. In noisy quantum computation, fresh qubits are known to be indispensable, since without them depolarizing noise permits only logarithmically deep computations [22–25]. How much memory a given supply of fresh qubits can buy has remained open.

Here we show that fresh qubits, not size, set a fundamental memory ceiling. Let k input qubits be reinitialized at each time step, and let every qubit undergo depolarizing noise that keeps its state with probability q. The largest memory depth that any such reservoir can attain, that is, the number of past inputs that still influence its state beyond an accuracy ε, grows exponentially with k and does not depend on the total number of qubits

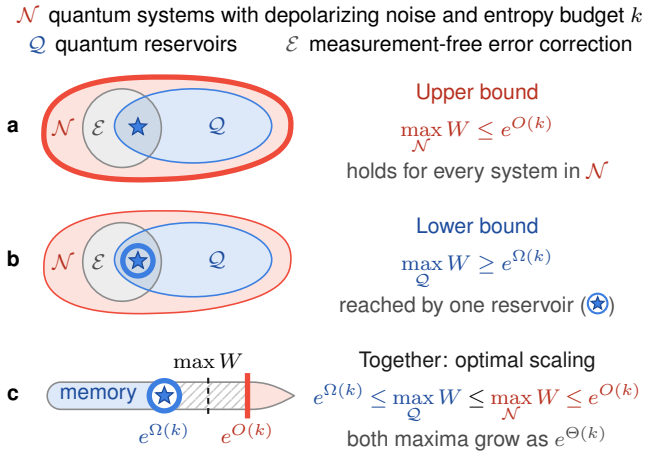


FIG. 1. **Scope of the memory limit** (schematic). (a) The upper bound of Theorem 1 holds for every system in $\mathcal{N}$, whatever its size and dynamics. (b) The lower bound is reached by one explicit reservoir (star): the error-correcting reservoir of Fig. 4(a), which corrects errors without measurements. (c) On the memory axis, no system goes beyond the upper bound and the star reaches the lower bound, so the memory ceiling $\max W$ lies in between and grows as $e^{\Theta(k)}$ [Fig. 3(b)].

n [Fig. 3(b)]:

$$W_{q,k}(\varepsilon) = \exp\big[\Theta_q(k)\big]. \tag{1}$$

This result is a fundamental resource limitation for an entire class of open quantum systems, not a property of a specific reservoir architecture [Fig. 1], and three features make it physically significant. First, it is a universal no-go result: within a well-defined noise model, no architecture, however large and including ones not yet conceived, beats the dimension-independent upper bound, and an explicit reservoir attains its exponential dependence on the refresh budget k. Second, it identifies entropy removal as the resource that limits memory: memory is bought with reinitialized qubits, not with size, and retaining an input for T time steps requires, and can be achieved with, $\Theta_q(\log T)$ reinitialized qubits per time step [Fig. 3(a)]. Third, it governs the stabilization of in-

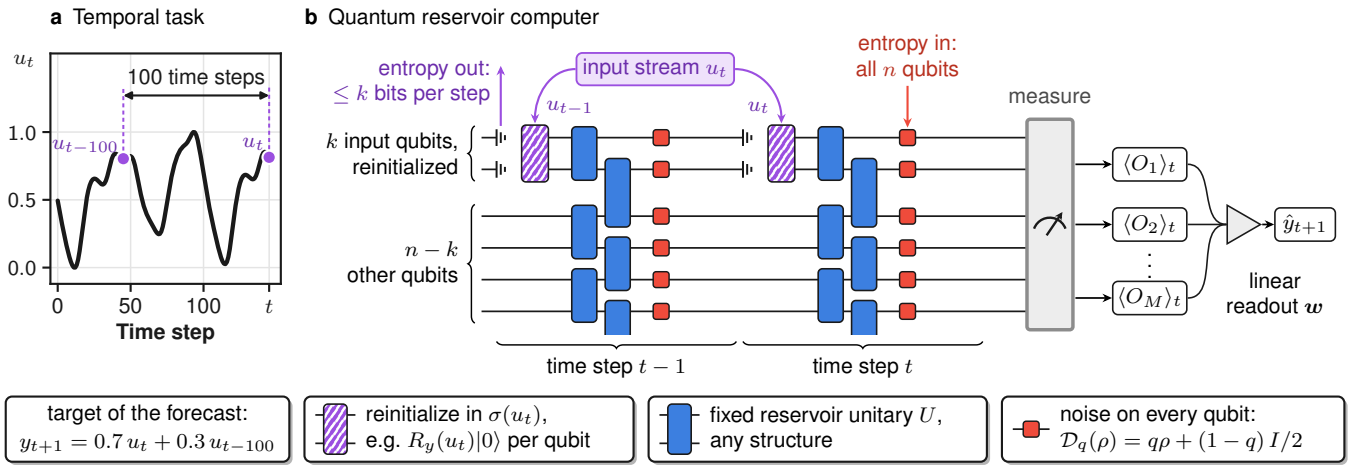


FIG. 2. **Noisy quantum reservoir computing.** (a) Forecasting $y_{t+1} = 0.7\,u_t + 0.3\,u_{t-100}$ needs a memory of 100 time steps. (b) Every time step reinitializes k input qubits in $\sigma(u_t)$, applies a fixed unitary U to all n qubits, and depolarizes every qubit; a linear readout predicts $\hat{y}_{t+1}$. Noise injects entropy into all n qubits, but only the k input qubits carry it out.

formation in noisy quantum systems in general: the proof uses only the entropy budget of each noisy time step, so the law applies to any noisy quantum dynamics refreshed by fresh qubits without classical feedback, from dynamical circuits with resets [11, 12] to measurement-free error correction [26], and complements space-overhead bounds for fault-tolerant quantum computation [27].

Our explicit error-correcting reservoir shows how fresh qubits sustain long memory under depolarizing noise. Its fixed dynamics protects the stored information while transferring accumulated entropy to the input register, where it is removed at the next reinitialization. The resulting memory depth grows exponentially with the number of reinitialized input qubits, realizing the optimal scaling established by our bounds.

A reservoir with fresh qubits.—A reservoir of n qubits receives one classical input u_t per time step t [Fig. 2(a),(b)]. At each time step, the k qubits of the input register are discarded and reinitialized in an input state $\sigma(u_t)$ that encodes the input. A fixed reservoir unitary U then acts on all qubits. Finally, every qubit undergoes depolarizing noise $\mathcal{D}_q(\rho) = q\rho + (1-q)\,\mathrm{Tr}(\rho)I/2$. Altogether,

$$\rho_t = \mathcal{D}_q^{\otimes n}\Big(U\big[\sigma(u_t) \otimes \mathrm{Tr}_{\mathrm{in}}\,\rho_{t-1}\big]U^\dagger\Big), \tag{2}$$

where $\mathrm{Tr}_{\mathrm{in}}$ is the partial trace over the input register. A trained linear readout turns the state into a prediction $\hat{y}_{t+1} = \sum_i w_i\,\mathrm{Tr}(O_i\rho_t) = \mathrm{Tr}(F\rho_t)$. This covers the original QRC proposal [6] and protocols that measure and reset qubits without feedback [11, 12]. We impose no restriction on $\sigma(u)$ or U; the unitary may even depend on the input.

The input qubits enter our analysis through a single number. Because $\sigma(u) \le I$, one time step can lower the entropy of a completely mixed reservoir by at most k bits; formally, the map $\mathcal{E}_u(X) = U[\sigma(u) \otimes \mathrm{Tr}_{\mathrm{in}}\,X]U^\dagger$ obeys $\mathcal{E}_u(I) \le 2^k I$. We call k the entropy *budget* of a time step. It is the only property of the input injection that the upper bound uses.

Noise erases the initial state, so the reservoir eventually reaches a state that depends only on the input history $h = (\ldots, u_{t-1}, u_t)$, its *echo state* $\rho(h)$ [1, 16, 28]. The *memory depth* $W(\varepsilon)$ is the smallest number L of time steps such that any two input histories that agree on their last L inputs lead to echo states with

$$\tfrac{1}{2}\|\rho(h) - \rho(h')\|_1 \le \varepsilon. \tag{3}$$

Inputs older than $W(\varepsilon)$ change every linear readout by at most $2\varepsilon\|F\|$ [28]: unlike the memory capacity [29], the memory depth limits all readouts at once. We write $W_{q,k}(\varepsilon)$ for the largest memory depth of any reservoir with k input qubits, maximized over n, U, and σ.

The memory limit.—Our main result bounds $W_{q,k}$ from both sides.

Theorem 1 (Memory limit)**.** *Let $0 < q < 1$ and $0 < \varepsilon < q$.*

(i) Every reservoir with $k \ge 4$ input qubits has

$$W(\varepsilon) \le \left\lceil 2^{k/(1-q)^2} \ln\frac{1}{\varepsilon} \right\rceil, \tag{4}$$

whatever its number of qubits, reservoir unitary, and input encoding.

(ii) For every k with $e^{\alpha_q(k-3)} \ge 4$, where $\alpha_q = \frac{1}{2}\ln\frac{1}{1-q^2}$, some reservoir with at most $2k$ qubits has

$$W(\varepsilon) \ge \tfrac{1}{4}\,e^{\alpha_q(k-3)} \ln\frac{q}{\varepsilon}. \tag{5}$$

Theorem 1 characterizes the optimal resource scaling over the entire class of admissible reservoirs, not the performance of a particular architecture: the upper bound

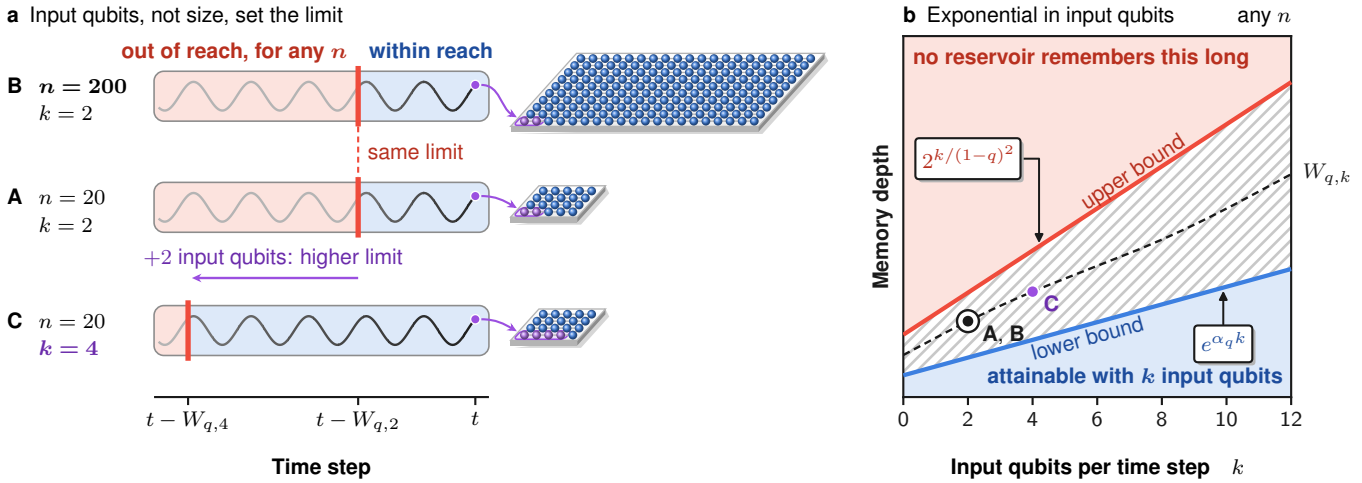


FIG. 3. **The memory limit** (schematic). (a) One input stream drives A and B, which differ tenfold in size but share $k = 2$ input qubits, and C, which is A with two more. Inputs older than the limit (red) are out of reach for every reservoir with the same k: A and B share the limit $W_{q,2}$, and C has a higher attainable limit. (b) The limit lies between two exponentials in k, for every n (Theorem 1); A, B, and C as in (a).

holds for every member, and the lower bound is an existence result, reached by one explicit member [Fig. 1]. Since the two exponents differ, the theorem fixes the scaling class, not the exact rate [Fig. 3(b)]. For $k \leq 3$, the upper bound holds with $k + 1$ in place of k [28].

The theorem concerns echo states: n redundant copies of an initial bit take $\Theta_q(\log n)$ time steps to wash out [28], so a larger reservoir needs a longer washout but does not remember new inputs for longer. Its constants are conservative [28]; it identifies the resource and its scaling.

Why no reservoir can remember longer.—A proof must defeat every strategy a reservoir might use to hide information: spreading it over many qubits, entangling it, or using its input qubits for repair. Our proof follows the difference between two reservoir states with different input histories and asks how well the states can be distinguished and how concentrated their difference is on a small part of the Hilbert space. Two established properties of depolarizing noise answer these questions [30, 31] and yield the bound in three steps.

Step 1: Noise caps the purity. We measure the purity of a state by $\log_2 P^2 = n - S_2(\rho)$, the number of pure qubits' worth of purity it holds, where S_2 is the Rényi-2 entropy. Depolarizing noise damps the many-qubit correlations of which purity is made, a property known as hypercontractivity [30]. As a result, every time step multiplies the log-purity by $\theta = 2q^2/(1 + q^2) < 1$, while the input qubits add at most θk. The log-purity thus has a fading memory of its own, and whatever its size, every echo state holds at most $\log_2 P_*^2 = 2kq^2/(1 - q^2)$ pure qubits' worth of purity, about 8.5 per input qubit at $q = 0.9$ [28].

Step 2: Spread-out differences are erased. Low purity alone does not mean forgetting: $|0\rangle\langle 0| \otimes I/2^{n-1}$ and $|1\rangle\langle 1| \otimes I/2^{n-1}$ are almost completely mixed, yet perfectly distinguishable. Their difference, however, is spread over the whole Hilbert space, and one layer of noise erases the fraction $1 - q$ of their distinguishability. A difference survives noise only if it is concentrated on a small subspace, as for a bit stored in many redundant copies. Reverse hypercontractivity makes this precise [31]: noise forces two states to overlap unless their difference is concentrated, so one layer of noise erases a fixed fraction of every spread-out difference.

Step 3: Concentration costs purity. A distinguishable difference concentrated on a fraction 2^{-m} of the Hilbert space, such as a bit stored in m redundant copies, requires about m pure qubits' worth of purity, and Step 1 caps the purity. A distance V that weighs distinguishability against concentration [28] therefore shrinks at every time step Φ_u, for any two reservoir states ρ, ρ' and any input u:

$$V\left(\Phi_u\rho - \Phi_u\rho'\right) \leq (1 - \delta)\, V(\rho - \rho'), \tag{6}$$

with a forgetting rate $\delta \geq 2^{-k/(1-q)^2}$ for $k \geq 4$ [28]. Whatever the reservoir does, every time step erases at least the fraction δ of its past. After L time steps, two echo states are therefore within trace distance $e^{-\delta L}$, which is Theorem 1(i).

Without fresh qubits, $k = 0$, as for inputs encoded in input-dependent unitaries, the echo state is the completely mixed state for every input sequence: such a reservoir retains nothing.

A reservoir that attains the limit.—Because $W_{q,k}$ is an optimum over all admissible reservoirs, a single reservoir suffices to show that the exponential scaling is attained. We deliberately choose the simplest one whose

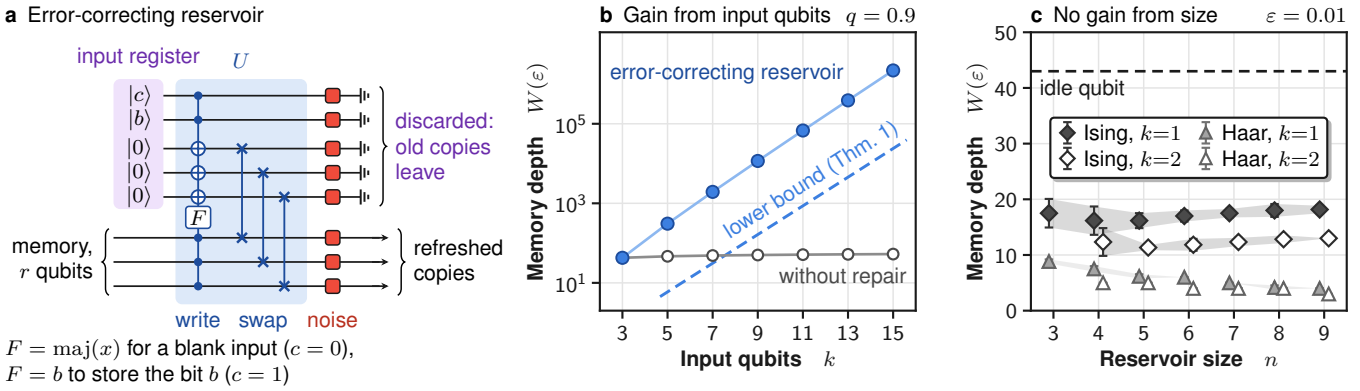


FIG. 4. **Attaining the limit.** (a) Error-correcting reservoir ($r = 3$, $k = 5$): U writes the new bit, or the majority of the stored copies, into fresh ancillas and swaps them into the memory; the old copies leave with the input register. (b) The error-correcting reservoir (exact single-qubit readout, $n = 2k - 2$) uses its fresh qubits for repair, and its memory depth grows exponentially with k, above the lower bound of Theorem 1; without repair, the same $r = k - 2$ copies gain almost nothing, even with the best readout. Depths with and without repair are plotted at continuous threshold crossings; the corresponding integer depths are rounded up. (c) Within the sampled ensembles, standard reservoirs gain nothing from size, and a second input qubit shortens their memory ($k = 1$ filled, $k = 2$ open; mean ± s.d., six realizations). Unitaries are sampled separately for each k. (b), (c) $q = 0.9$, $\varepsilon = 0.01$.

dynamics is exactly solvable; its role is to show that the exponential dependence on k is fundamental rather than an artifact of the upper-bound proof.

The idea is a memory that repairs itself with fresh qubits, an error-correcting reservoir [Fig. 4(a)]. It stores a bit in r memory qubits as a repetition code and corrects it by majority vote: at every time step, its unitary computes the majority of the stored copies, writes it into r freshly reinitialized ancillas, and swaps the ancillas into the memory. The old copies, which carry the noise of the previous time step, end up in the input register and are discarded at the next injection, so the entropy leaves with them. Two further input qubits tell the reservoir whether to keep its memory or to store a new bit, so that $k = r + 2$ [28]. The unitary only permutes computational basis states, so the reservoir corrects errors without measurements or feedback.

Because all states remain diagonal, every qubit behaves as a classical bit that flips with probability $(1 - q)/2$ per time step, and the response can be computed exactly. The stored bit is lost only if most of its r copies flip within one time step, which happens with probability $p_r \le (1 - q^2)^{r/2} = e^{-\alpha_q r}$ by a Chernoff bound with the exact exponent [28]. A bit b stored t time steps ago therefore yields on any memory qubit

$$\langle Z \rangle_t = (-1)^b \, q \, (1 - 2p_r)^t. \tag{7}$$

The signal still decays, but at a rate that is exponentially small in r. Two input histories that differ only in this bit agree on their last t inputs, yet their echo states remain at trace distance at least $q(1 - 2p_r)^t$; with the bound on p_r, this gives Theorem 1(ii). At $q = 0.9$ and $\varepsilon = 0.01$, the memory depth exceeds 300 time steps for $k = 5$ and two million for $k = 15$, whereas the same copies without the majority vote keep the bit for at most 53 time steps [Fig. 4(b)].

Standard reservoirs operate far below the limit.—We simulated transverse-field Ising reservoirs with random all-to-all couplings, the workhorse of QRC [6, 7, 32], and Haar-random unitaries at $q = 0.9$, and estimated their memory depth W at $\varepsilon = 0.01$ from input histories that differ in one input or in all earlier inputs [28]. Within these ensembles, size does not extend their memory [Fig. 4(c)]: with one input qubit ($k = 1$), the Ising reservoir remembers 16 to 18 time steps for every size from $n = 3$ to 9 qubits, and the Haar reservoir forgets faster as it grows (W drops from 9 to 4 time steps as n grows from 3 to 9). More input qubits do not help either [Fig. 4(c)] [28]: at fixed size ($n = 8$), raising k from 1 to 4 lowers W from 18 to 8 time steps (Ising) and from 4 to 3 (Haar); the decrease persists when the additional input qubits are reinitialized in $|0\rangle$. Both stay below an idle qubit that merely holds the input ($W = 43$ time steps), whereas the error-correcting reservoir of the same size ($n = 8$, of which $k = 5$ are input qubits) remembers 316 time steps.

An entropy budget.—Beyond depolarizing noise, the result becomes an entropy budget [28]: any additional channel $\mathcal{S}$ with $\mathcal{S}(I) \le 2^b I$ adds b bits, and the upper bound holds with k replaced by $k + b$. Unital noise (Pauli errors, dephasing, coherent errors, crosstalk) and projective measurements whose outcomes only enter the readout [33] have $b = 0$; since every unital qubit channel with full-rank Choi matrix contains a depolarizing component, the theorem covers generic independent unital noise [28]. Resets of m qubits count like m input

qubits. Amplitude damping with probability $0 < \gamma < 1$ adds $n \log_2(1+\gamma)$ bits, and relaxing idle qubits then act as a refrigerator [23] that lets a reservoir with two input qubits remember for a time exponential in n [28, 34]. Without a depolarizing component, with non-Markovian noise [18], or with feedback of measurement outcomes, memory can be unbounded [28].

Conclusion.—Fresh qubits, not size, set the memory of noisy quantum reservoirs. A dimension-independent upper bound for arbitrary dynamics and an explicit reservoir that attains its exponential scaling establish an optimal resource law and a universal no-go result for an entire class of open quantum systems, from QRC to measurement-free error correction [26]: with k reinitialized qubits per time step, the memory ceiling is $\exp[\Theta_q(k)]$ time steps, whatever the size. Fresh qubits are necessary to reach the ceiling but not sufficient; standard reservoirs, which do not use them for repair, operate far below it. When a reservoir forgets too quickly, the parameters to check are therefore its supply of fresh qubits and whether its dynamics uses them to remove errors, not its size. The theorem bounds how far back any task can reach, not how well a reservoir processes the inputs within that range [35]. Open problems include closing the gap between the two exponents, extending the bound to feedback and non-Markovian noise, and testing the scaling by varying k at fixed n with mid-circuit resets [11, 12].

Acknowledgments.—Markus Baumann used Claude Opus 5.5 (Anthropic) to assist with drafting and editing the text, checking derivations and references, and writing the simulation and plotting code. Markus Baumann directed this work, verified every derivation and numerical result, and takes full responsibility for the content of this article.

Data availability.—Code and data that reproduce all numbers and figures of this Letter are provided as a source archive with the Supplemental Material [28].

* Contact author: markus.baumann@campus.lmu.de

[1] H. Jaeger, *The "echo state" approach to analysing and training recurrent neural networks*, GMD Report 148 (German National Research Center for Information Technology, 2001).
[2] W. Maass, T. Natschläger, and H. Markram, Real-time computing without stable states: A new framework for neural computation based on perturbations, Neural Comput. **14**, 2531 (2002).
[3] H. Jaeger and H. Haas, Harnessing nonlinearity: Predicting chaotic systems and saving energy in wireless communication, Science **304**, 78 (2004).
[4] M. Lukoševičius and H. Jaeger, Reservoir computing approaches to recurrent neural network training, Comput. Sci. Rev. **3**, 127 (2009).
[5] G. Tanaka, T. Yamane, J. B. Héroux, R. Nakane, N. Kanazawa, S. Takeda, H. Numata, D. Nakano, and A. Hirose, Recent advances in physical reservoir computing: A review, Neural Netw. **115**, 100 (2019).
[6] K. Fujii and K. Nakajima, Harnessing disordered-ensemble quantum dynamics for machine learning, Phys. Rev. Appl. **8**, 024030 (2017).
[7] K. Nakajima, K. Fujii, M. Negoro, K. Mitarai, and M. Kitagawa, Boosting computational power through spatial multiplexing in quantum reservoir computing, Phys. Rev. Appl. **11**, 034021 (2019).
[8] P. Mujal, R. Martínez-Peña, J. Nokkala, J. García-Beni, G. L. Giorgi, M. C. Soriano, and R. Zambrini, Opportunities in quantum reservoir computing and extreme learning machines, Adv. Quantum Technol. **4**, 2100027 (2021).
[9] J. Chen, H. I. Nurdin, and N. Yamamoto, Temporal information processing on noisy quantum computers, Phys. Rev. Appl. **14**, 024065 (2020).
[10] Y. Suzuki, Q. Gao, K. C. Pradel, K. Yasuoka, and N. Yamamoto, Natural quantum reservoir computing for temporal information processing, Sci. Rep. **12**, 1353 (2022).
[11] T. Yasuda, Y. Suzuki, T. Kubota, K. Nakajima, Q. Gao, W. Zhang, S. Shimono, A. Matsuo, J. Ananthanarayanan, H. I. Nurdin, and N. Yamamoto, Quantum reservoir computing with repeated measurements on superconducting devices (2023), arXiv:2310.06706 [quant-ph].
[12] F. Hu, S. A. Khan, N. T. Bronn, G. Angelatos, G. E. Rowlands, G. J. Ribeill, and H. E. Türeci, Overcoming the coherence time barrier in quantum machine learning on temporal data, Nat. Commun. **15**, 7491 (2024).
[13] F. Monzani, E. Ricci, L. Nigro, and E. Prati, Non-unital noise in a superconducting quantum computer as a computational resource for reservoir computing (2024), arXiv:2409.07886 [quant-ph].
[14] T. Dao, E. Yilmaz, I. Shehzad, C. Pere, K. Ghosh, I. Wittmann, T. Brunschwiler, G. Cortiana, C. O'Meara, S. Woerner, and F. Tacchino, Breaking concentration barriers for quantum extreme learning on digital quantum processors (2026), arXiv:2603.13005 [quant-ph].
[15] S. Boyd and L. O. Chua, Fading memory and the problem of approximating nonlinear operators with Volterra series, IEEE Trans. Circuits Syst. **32**, 1150 (1985).
[16] R. Martínez-Peña and J.-P. Ortega, Quantum reservoir computing in finite dimensions, Phys. Rev. E **107**, 035306 (2023).
[17] O. Morguí-Sancho, G. Manzano, G. L. Giorgi, and R. Zambrini, General theory of monitored quantum reservoir computing (2026), arXiv:2607.14311 [quant-ph].
[18] A. Sannia, R. Ravell Rodríguez, G. L. Giorgi, and R. Zambrini, Non-Markovianity and memory enhancement in quantum reservoir computing, npj Quantum Inf. **12**, 111 (2026).
[19] T. Kubota, Y. Suzuki, S. Kobayashi, Q. H. Tran, N. Yamamoto, and K. Nakajima, Temporal information processing induced by quantum noise, Phys. Rev. Res. **5**, 023057 (2023).
[20] A. Sannia, R. Martínez-Peña, M. C. Soriano, G. L. Giorgi, and R. Zambrini, Dissipation as a resource for quantum reservoir computing, Quantum **8**, 1291 (2024).
[21] J. Freiheit and F. Campaioli, Superextensive learning in quantum reservoirs at the onset of information scrambling (2026), arXiv:2608.25511 [quant-ph].
[22] D. Aharonov, M. Ben-Or, R. Impagliazzo, and N. Nisan, Limitations of noisy reversible computation (1996), arXiv:quant-ph/9611028.

[23] M. Ben-Or, D. Gottesman, and A. Hassidim, Quantum refrigerator (2013), arXiv:1301.1995 [quant-ph].

[24] D. Stilck França and R. García-Patrón, Limitations of optimization algorithms on noisy quantum devices, Nat. Phys. **17**, 1221 (2021).

[25] A. A. Mele, A. Angrisani, S. Ghosh, S. Khatri, J. Eisert, D. Stilck França, and Y. Quek, Noise-induced shallow circuits and the absence of barren plateaus, Nat. Phys. **22**, 751 (2026).

[26] S. Heußen, D. F. Locher, and M. Müller, Measurement-free fault-tolerant quantum error correction in near-term devices, PRX Quantum **5**, 010333 (2024).

[27] O. Fawzi, A. Müller-Hermes, and A. Shayeghi, A lower bound on the space overhead of fault-tolerant quantum computation, in *13th Innovations in Theoretical Computer Science Conference (ITCS 2022)*, Leibniz International Proceedings in Informatics (LIPIcs), Vol. 215 (Schloss Dagstuhl–Leibniz-Zentrum für Informatik, 2022) pp. 68:1–68:20.

[28] See the Supplemental Material accompanying this manuscript for the complete proofs, the exact analysis of the error-correcting reservoir, the entropy budget beyond depolarizing noise, the details of the simulations, and the Python code that reproduces all numerical values and data figures, which includes Refs. [36–38]. Code repository: `https://github.com/eybmits/fresh-qubits-memory`.

[29] H. Jaeger, *Short term memory in echo state networks*, GMD Report 152 (German National Research Center for Information Technology, 2002).

[30] C. King, Hypercontractivity for semigroups of unital qubit channels, Commun. Math. Phys. **328**, 285 (2014).

[31] T. Cubitt, M. Kastoryano, A. Montanaro, and K. Temme, Quantum reverse hypercontractivity, J. Math. Phys. **56**, 102204 (2015).

[32] R. Martínez-Peña, G. L. Giorgi, J. Nokkala, M. C. Soriano, and R. Zambrini, Dynamical phase transitions in quantum reservoir computing, Phys. Rev. Lett. **127**, 100502 (2021).

[33] P. Mujal, R. Martínez-Peña, G. L. Giorgi, M. C. Soriano, and R. Zambrini, Time-series quantum reservoir computing with weak and projective measurements, npj Quantum Inf. **9**, 16 (2023).

[34] W. Hoeffding, Probability inequalities for sums of bounded random variables, J. Am. Stat. Assoc. **58**, 13 (1963).

[35] J. Cao, R.-Y. Gong, Z. Lin, Y. Zeng, and Z.-L. Xiang, Liouvillian response for temporal information in quantum reservoir computing (2026), arXiv:2609.23429 [quant-ph].

[36] J. von Neumann, Probabilistic logics and the synthesis of reliable organisms from unreliable components, in *Automata Studies*, edited by C. E. Shannon and J. McCarthy (Princeton University Press, Princeton, 1956) p. 43.

[37] J. Nelson, J. Rajakumar, and M. J. Gullans, Non-Clifford gates are required for long-term memory, in *41st Computational Complexity Conference (CCC 2026)*, Leibniz International Proceedings in Informatics (LIPIcs), Vol. 383 (Schloss Dagstuhl–Leibniz-Zentrum für Informatik, 2026) pp. 38:1–38:13.

[38] W. Xiong, G. Facelli, M. Sahebi, O. Agnel, T. Chotibut, S. Thanasilp, and Z. Holmes, On fundamental aspects of quantum extreme learning machines, Quantum Mach. Intell. **7**, 20 (2025).

Supplemental Material: The Fundamental Memory Ceiling of Noisy Quantum Reservoirs

Markus Baumann, Jonas Stein, and Claudia Linnhoff-Popien

Quantum Applications and Research Laboratory (QAR-Lab),
Chair of Mobile and Distributed Systems, Department of Computer Science,
Ludwig-Maximilians-Universität München, Oettingenstraße 67, 80538 München, Germany

This Supplemental Material contains the proofs of the main text, the exact analysis of the error-correcting reservoir, the entropy budget for noise beyond depolarizing noise, the details of the simulations, and a guide to the code and data. Section I explains the idea of the proof in words, and every later section can be read on its own. Equations, figures, tables, and lemmas carry the prefix S; numbers without prefix and references refer to the main text.

CONTENTS

I. How the proof works S2

II. What the proof keeps track of S2
A. Averaged norms S2
B. Distinguishability, purity, and concentration S2
C. The fresh qubits are the only source of purity S3

III. Proof of the upper bound S3
A. Two properties of depolarizing noise S3
B. Step 1: Noise caps the purity S3
C. Step 2: Spread-out differences are erased S4
D. Step 3: One distance shrinks at every time step S5
E. How large is the forgetting rate? S5
F. From a shrinking distance to the memory bound S6
G. How tight are the bounds? S6

IV. The washout grows with the reservoir size S7

V. The error-correcting reservoir S8
A. The construction S8
B. Exact response S8
C. The majority error and the exponent α_q S8
D. Proof of Theorem 1(ii) S9
E. Further properties S9

VI. Entropy budget beyond depolarizing noise S10
A. Unital noise and measurements S10
B. Non-unital noise and resets S11
C. Every unital qubit channel with full-rank Choi matrix contains depolarizing noise S11
D. Cooling noise: memory exponential in n S11
E. Where the limit ends: two counterexamples S13

VII. Numerical simulations S13
A. Reservoirs S13
B. Memory depth S13
C. Validation S14
D. Results and robustness S15

VIII. Implications for reservoir design, and scope S15
A. Design lessons S15
B. What the result does not say S16

IX. Code and data S16

I. HOW THE PROOF WORKS

The upper bound of Theorem 1 rests on a single statement: every time step shrinks the difference between any two reservoir states by a fixed fraction, whatever the reservoir does. The difficulty is that neither of the two obvious ways to measure a difference shrinks on its own.

- The *trace distance*, which says how well two states can be distinguished, can barely change in one time step: a bit stored in many redundant copies passes a layer of noise almost intact.
- The *2-norm*, which is large when a difference is concentrated on a small part of the Hilbert space, can even grow, because the fresh input qubits make the state purer.

The proof combines the two in the three steps of the main text.

1. *Noise caps the purity* of every state (Lemma S2).
2. *One layer of noise erases a fixed fraction of every difference that is not concentrated* (Lemma S3).
3. *Concentration costs purity*, so a combination of trace distance and 2-norm shrinks at every time step (Proposition S4).

Steps 1 and 2 each rest on one known inequality for depolarizing noise, hypercontractivity [30] and reverse hypercontractivity [31]; all other steps are elementary. Section III F turns the shrinking distance into the memory bound of Theorem 1(i). Sections IV–VII derive the remaining statements of the main text: the washout of the initial state, the error-correcting reservoir and Theorem 1(ii), the entropy budget, and the simulations.

II. WHAT THE PROOF KEEPS TRACK OF

A. Averaged norms

For an operator X on n qubits, let $d = 2^n$ and $\tau(X) = \mathrm{Tr}(X)/d$, the average eigenvalue of X. We use

$$|X|_p = \big(\tau|X|^p\big)^{1/p}, \tag{S1}$$

the p-th root of the average p-th power of the singular values of X (the normalized Schatten p-norm). Averaging instead of summing makes $|I|_p = 1$ for every p, so that no factors of d appear in the proofs. We need $p = 1$, $p = 2$, and, in one place, a value of p between 0 and 1, for which the same formula defines a useful quantity, although not a norm. Three standard facts about such averages are used: (i) $|X|_p$ grows with p; (ii) $p \mapsto \ln \tau|X|^p$ is convex; and (iii) $|X|_p \le |X|_1^{1-\theta}|X|_2^{\theta}$ whenever $1/p = 1 - \theta/2$, which is Hölder's inequality.

B. Distinguishability, purity, and concentration

The proof follows the operator $X = d(\rho - \rho')$ for two reservoir states ρ and ρ'. Three properties of it matter.

- *Distinguishability*: $|X|_1 = \|\rho - \rho'\|_1$ is twice the trace distance. The best measurement tells ρ from ρ' with success probability $\frac{1}{2} + \frac{1}{4}|X|_1$.
- *Purity*: for a single state, $P = |d\rho|_2 = (d\,\mathrm{Tr}\,\rho^2)^{1/2}$, and $\log_2 P^2 = n - S_2(\rho)$, where $S_2 = -\log_2 \mathrm{Tr}\,\rho^2$ is the Rényi-2 entropy. It counts *pure qubits' worth* of purity: $\log_2 P^2$ equals 0 for the completely mixed state, n for a pure state, and m for a state with m pure and $n - m$ completely mixed qubits.
- *Concentration*: the ratio $|X|_2/|X|_1 \ge 1$. It equals one if all eigenvalues of X have the same magnitude, and it is large if X lives on a small subspace.

Three examples show the scale of the concentration. The difference between $|0\rangle\langle 0| \otimes I/2^{n-1}$ and $|1\rangle\langle 1| \otimes I/2^{n-1}$ has concentration 1: it is spread over the whole Hilbert space. A bit stored in m redundant copies, $|0\rangle\langle 0|^{\otimes m}$ versus $|1\rangle\langle 1|^{\otimes m}$ with the other qubits completely mixed, has concentration $2^{(m-1)/2}$. Two orthogonal pure states have the largest possible concentration, $2^{(n-1)/2}$. In all three examples, a concentrated difference requires purity, and this holds in general: since $|X|_2 \le |d\rho|_2 + |d\rho'|_2$, two states with purity at most P have a difference with $|X|_2 \le 2P$.

C. The fresh qubits are the only source of purity

Lemma S1 (Budget). *The injection step $\mathcal{E}_u(X) = U[\sigma(u) \otimes \mathrm{Tr}_{\mathrm{in}} X]U^\dagger$ obeys, for every Hermitian X,*

$$|\mathcal{E}_u X|_1 \le |X|_1, \qquad |\mathcal{E}_u X|_2 \le 2^{k/2}\,|X|_2. \tag{S2}$$

Proof. The first inequality says that no physical process makes two states more distinguishable (contractivity of the trace norm). For the second, write $X = \sum_{i,j} |i\rangle\langle j| \otimes X_{ij}$, where $|i\rangle$ runs over the 2^k basis states of the input register. Then $\mathrm{Tr}_{\mathrm{in}} X = \sum_i X_{ii}$, and the Cauchy–Schwarz inequality gives $\mathrm{Tr}(\mathrm{Tr}_{\mathrm{in}} X)^2 \le 2^k \sum_i \mathrm{Tr}\, X_{ii}^2 \le 2^k \,\mathrm{Tr}\, X^2$. Since $\mathrm{Tr}\,\sigma(u)^2 \le 1$ and U does not change traces, $\mathrm{Tr}(\mathcal{E}_u X)^2 = \mathrm{Tr}\,\sigma(u)^2\ \mathrm{Tr}(\mathrm{Tr}_{\mathrm{in}} X)^2 \le 2^k \,\mathrm{Tr}\, X^2$. Dividing by d gives the claim. □

In words, discarding k qubits raises the 2-norm of an operator by at most a factor $2^{k/2}$. For a state, $X = d\rho$, a time step thus adds at most k pure qubits' worth of purity: the fresh qubits are the only source of purity.

The same two bounds hold, with k replaced by b, for every channel $\mathcal{E}$ with $\mathcal{E}(I) \le 2^b I$; for the 2-norm, this follows from the Kadison–Schwarz inequality $(\mathcal{E}X)^2 \le \|\mathcal{E}(I)\|\,\mathcal{E}(X^2)$. We call b the *budget* of $\mathcal{E}$. The injection step has $b = k$, because $\mathcal{E}_u(I) = 2^k U[\sigma(u) \otimes I]U^\dagger \le 2^k I$; unitaries and all unital channels have $b = 0$; and a reset of m qubits has $b = m$ (Sec. VI). In the following, $\mathcal{E}$ denotes any channel with a budget of b bits, and $b = k$ for the reservoirs of the main text.

III. PROOF OF THE UPPER BOUND

Throughout this section, $\mathcal{N}_q = \mathcal{D}_q^{\otimes n}$ denotes one layer of depolarizing noise, so that one time step is $\Phi_u = \mathcal{N}_q \circ \mathcal{E}_u$.

A. Two properties of depolarizing noise

Noise smooths concentrated operators (hypercontractivity). For every Hermitian Y,

$$|\mathcal{N}_q Y|_2 \le |Y|_{1+q^2}. \tag{S3}$$

This is King's hypercontractivity theorem for products of unital qubit channels [30, Corollary 2], which gives $|\mathcal{N}_q Y|_r \le |Y|_p$ whenever $q^2 \le (p-1)/(r-1)$; here $r = 2$ and $p = 1 + q^2$. Because $1 + q^2 < 2$, the norm on the right is weaker than the one on the left, so noise turns operators with a bounded $(1+q^2)$-norm into operators with a bounded 2-norm. The mechanism is the damping of correlations: written as a sum of Pauli strings, an operator loses the factor q for every qubit on which a string acts, and concentrated operators, such as projectors onto a few states, consist mostly of long strings. As a single-qubit check, $Y = 2|0\rangle\langle 0| = I + Z$ becomes $I + qZ$, and Eq. (S3) reads $1.345 \le 1.364$ at $q = 0.9$. (The extension from positive to Hermitian $Y = Y_+ - Y_-$ is immediate: $\mathcal{N}_q$ is positive and self-adjoint with respect to τ, so $|\mathcal{N}_q Y|_2^2 = \tau[Y\mathcal{N}_q^2(Y)] \le \tau[|Y|\mathcal{N}_q^2(|Y|)] = |\mathcal{N}_q|Y||_2^2$, while $||Y||_p = |Y|_p$.)

Noise forces overlap (reverse hypercontractivity). For positive operators $G_1, G_2 \ge 0$,

$$\tau\big[G_1\,\mathcal{N}_q(G_2)\big] \ge |G_1|_{1-q}\,|G_2|_{1-q}. \tag{S4}$$

This is the two-point reverse hypercontractive inequality of Cubitt *et al.* [31, Corollary 14], which holds for all exponents $\bar{p}_1, \bar{p}_2 < 1$ with $q \le \sqrt{(1-\bar{p}_1)(1-\bar{p}_2)}$; we use the symmetric choice $\bar{p}_1 = \bar{p}_2 = 1 - q$. To read it, normalize $\tau G = 1$, so that G is d times a density matrix. Then $|G|_{1-q}$ lies between 0 and 1: it equals one for $G = I$ and is small only if G is concentrated on a small subspace. After one layer of noise, two positive operators therefore overlap substantially unless they are concentrated. For example, the scaled orthogonal projectors $2|0\rangle\langle 0|$ and $2|1\rangle\langle 1|$ of a single qubit overlap by $1 - q$ after one layer of noise.

B. Step 1: Noise caps the purity

Lemma S2 (Purity cap). *Let $\theta = 2q^2/(1+q^2)$. For every Hermitian X,*

$$|\mathcal{N}_q \mathcal{E} X|_2 \le 2^{b\theta/2}\,|X|_1^{1-\theta}\,|X|_2^{\theta}. \tag{S5}$$

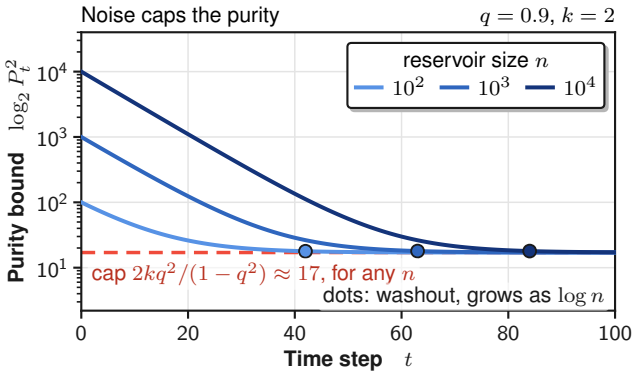


FIG. S1. **Noise caps the purity.** Purity bound $\log_2 P_t^2$ obtained by iterating Eq. (S6) with equality, for reservoirs of n qubits that start in a pure state, at $q = 0.9$ and $b = k = 2$. Every bound approaches the same cap $\log_2 P_*^2 = 2kq^2/(1-q^2)$ (dashed), whatever the size; only the washout time t_0 of Sec. III F (dots) grows, as $\log n$.

Proof. Hypercontractivity (S3), fact (iii) with $p = 1 + q^2$ (which fixes θ), and Lemma S1 give

$$|\mathcal{N}_q \mathcal{E} X|_2 \le |\mathcal{E} X|_{1+q^2} \le |\mathcal{E} X|_1^{1-\theta} \, |\mathcal{E} X|_2^{\theta} \le |X|_1^{1-\theta} \big(2^{b/2} |X|_2\big)^{\theta}. \qquad \square$$

After the noise, the 2-norm is thus controlled by a weighted geometric mean of the 1-norm and the 2-norm before it. For a state, $X = d\rho$ has $|X|_1 = 1$, so the purity obeys $P_t \le 2^{b\theta/2} P_{t-1}^{\theta}$, or

$$\log_2 P_t \le \theta \log_2 P_{t-1} + \tfrac{1}{2} b\theta. \tag{S6}$$

The log-purity is multiplied by $\theta < 1$ at every time step, and the budget adds at most $b\theta/2$. Its fixed point is the *purity cap*

$$\log_2 P_* = \frac{bq^2}{1-q^2}, \tag{S7}$$

which is the purity cap of the main text for $b = k$. Whatever the size of the reservoir, the purity approaches the same cap; only the time it takes grows, as $\log n$ [Fig. S1 and Sec. III F].

In the proofs below, we use the cap in the form $2^{b\theta/2} = P_*^{\bar{\theta}}$ with $\bar{\theta} = 1 - \theta = (1-q^2)/(1+q^2)$.

C. Step 2: Spread-out differences are erased

Lemma S3 (Forced forgetting)**.** *Let* $a = 2q/(1-q)$ *and* $c_q = 2^{-1/(1-q)}$. *Every traceless Hermitian* $Y \neq 0$ *with concentration* $K = |Y|_2/|Y|_1$ *obeys*

$$|\mathcal{N}_q Y|_1 \le |Y|_1 \big[1 - c_q \, (2K)^{-a}\big]. \tag{S8}$$

One layer of noise thus removes at least the fraction $c_q(2K)^{-a}$ of the distinguishability. This fraction is small only if the difference is concentrated.

Proof. Split the difference. Write $Y = Y_+ - Y_-$ with $Y_\pm \ge 0$ supported on orthogonal subspaces. Because Y is traceless, both parts have the same weight, $\tau Y_\pm = \frac{1}{2}|Y|_1$, so $G_\pm = 2Y_\pm/|Y|_1$ are d times density matrices. Their purity is limited by the concentration of Y: $|G_\pm|_2 \le 2K$.

Distinguishability is the success of the best test. A measurement with outcomes M and $I - M$ that guesses "+" on M confuses the noisy parts with error probabilities that sum to $\mathrm{err}(M) = \tau[(I-M)\mathcal{N}_q G_+] + \tau[M\mathcal{N}_q G_-]$. The trace norm is attained by the best such test, which gives

$$|\mathcal{N}_q Y|_1 = |Y|_1 \Big[1 - \min_{0 \le M \le I} \mathrm{err}(M)\Big]. \tag{S9}$$

Every test errs. Either M or $I - M$ accepts at least half of the Hilbert space; say $\tau M \ge \frac{1}{2}$. Reverse hypercontractivity (S4) then bounds the probability that M accepts the noisy negative part: $\mathrm{err}(M) \ge \tau[M\mathcal{N}_q G_-] \ge |M|_{1-q} \, |G_-|_{1-q}$.

Both factors are bounded from below. A test that accepts half of the Hilbert space is not concentrated: $M^{1-q} \geq M$ gives $|M|_{1-q} \geq (\tau M)^{1/(1-q)} \geq c_q$. A state part with limited purity is not concentrated either: by fact (ii), $\ln \tau G_-^p$ is convex in p, and because $\tau G_- = 1$ at $p = 1$, which lies between $1-q$ and 2, this gives $\tau G_-^{1-q} \geq (\tau G_-^2)^{-q}$, that is, $|G_-|_{1-q} \geq |G_-|_2^{-a} \geq (2K)^{-a}$. Hence $\mathrm{err}(M) \geq c_q (2K)^{-a}$ for every test, and Eq. (S9) proves the lemma. □

D. Step 3: One distance shrinks at every time step

Lemma S3 needs a bound on the concentration, and Lemma S2 provides one. We therefore measure a difference by both norms at once,

$$V(X) = \max\Big\{|X|_1,\ \frac{|X|_2}{2P_*}\Big\}. \tag{S10}$$

In words, V is twice the trace distance, unless the difference is so concentrated that its 2-norm, measured in units of twice the purity cap, is larger. For the difference of two states with purity at most P_*, both entries are at most 2 (Sec. II).

Proposition S4 (Contraction). *For every traceless Hermitian X,*

$$V(\mathcal{N}_q \mathcal{E} X) \leq (1-\delta)\, V(X), \qquad \delta = c_q \big(4 \cdot 2^{b/2} P_*\big)^{-a}. \tag{S11}$$

Proof. Both sides grow linearly with X, so it suffices to take $V(X) = 1$, that is, $|X|_1 \leq 1$ and $|X|_2 \leq 2P_*$.

The 2-norm part shrinks. Lemma S2 and $2^{b\theta/2} = P_*^{\bar\theta}$ give $|\mathcal{N}_q \mathcal{E} X|_2 \leq P_*^{\bar\theta} (2P_*)^{\theta} = 2^{-\bar\theta} \cdot 2P_*$, so the 2-norm part of V shrinks by the factor $2^{-\bar\theta}$.

The 1-norm part shrinks. The operator $Y = \mathcal{E}X$ is traceless, with $x = |Y|_1 \leq 1$ and, by Lemma S1, $|Y|_2 \leq 2^{b/2} \cdot 2P_*$. If $Y \neq 0$, its concentration is therefore at most $2^{b/2} \cdot 2P_*/x$, and Lemma S3 gives $|\mathcal{N}_q Y|_1 \leq x - \delta x^{a+1}$; for $Y = 0$, this holds trivially. The right-hand side increases with x on $[0,1]$, because $(a+1)\delta \leq 1$, so $|\mathcal{N}_q Y|_1 \leq 1 - \delta$.

Both parts are at most $1-\delta$, because $2^{-\bar\theta} \leq 1 - \delta$. □

The two small facts used at the end, $(a+1)\delta \leq 1$ and $\delta \leq 1 - 2^{-\bar\theta}$, follow from one elementary inequality. With $s = 1/(1-q) \geq 1$, we have $\delta \leq c_q 4^{-a} = 2^{4-5s}$, $a + 1 \leq 2s$, and $1 - 2^{-\bar\theta} \geq \bar\theta/2 \geq (1-q)/2 = 1/(2s)$. Both facts thus follow from $2s \cdot 2^{4-5s} \leq 1$, that is, $\log_2 s \leq 5(s-1)$, which holds for all $s \geq 1$. No case distinction and no strong-noise condition are needed, and the proposition holds for every channel with a budget of b bits, including input-dependent unitaries.

E. How large is the forgetting rate?

Corollary S5 (Rate). *Inserting c_q, a, and P_* gives*

$$\delta^{-1} = 2^{(1+4q)/(1-q)} \cdot 2^{\beta_q b}, \qquad \beta_q = \frac{q(1+q^2)}{(1-q)(1-q^2)}. \tag{S12}$$

In particular, $\delta^{-1} \leq 2^{b/(1-q)^2}$ for $b \geq 4$, and $\delta^{-1} \leq 2^{(b+1)/(1-q)^2}$ for $b \geq 1$.

Proof. The four factors of $\delta^{-1} = c_q^{-1}\, 4^a\, 2^{ab/2}\, P_*^a$ contribute $\frac{1}{1-q}$, $\frac{4q}{1-q}$, $\frac{qb}{1-q}$, and $\frac{2q^3 b}{(1-q)(1-q^2)}$ to $\log_2 \delta^{-1}$; the first two add up to $(1+4q)/(1-q)$, the last two to $\beta_q b$. For the simple forms, compare exponents after multiplying by $(1-q)^2$, using $(1-q)^2 \beta_q = q(1+q^2)/(1+q)$. The first form becomes $b(1+q+q^2) \geq (1+4q)(1+q)$, which holds for $b \geq 4$ because $4(1+q+q^2) - (1+4q)(1+q) = 3 - q > 0$. The second becomes $b(1-q^3) + q(4q-3)(1+q) \geq 0$, which holds for $b \geq 1$ because $1 - 3q + q^2 + 3q^3 > 0$ for $0 \leq q \leq 1$. □

The inverse forgetting rate δ^{-1} grows by the factor 2^{β_q} for every bit of budget. At $q = 0.9$, the constants are $\theta = 0.895$, $a = 18$, $c_q = 2^{-10}$, and $\beta_q = 85.7$, so $\delta^{-1} = 2^{85.7b+46}$, compared with 2^{100b} from the first simple form. For $b = 3$, the first simple form fails for $q > \sqrt{3} - 1 \approx 0.73$, which is why Theorem 1(i) assumes $k \geq 4$.

F. From a shrinking distance to the memory bound

Washout of the purity. An arbitrary initial state has $\log_2 P_0 \le n/2$. Iterating Eq. (S6) gives

$$\log_2 P_t \le \log_2 P_* + \theta^t \frac{n}{2}. \tag{S13}$$

The excess purity of the initial state thus decays geometrically toward the cap. After $t_0 = \lceil \ln n/(-\ln\theta)\rceil$ time steps, $\theta^t n \le 1$, so the initial state contributes at most one pure qubit's worth, $P_t \le \sqrt{2}P_*$, and one more time step keeps it there, because $2^{b\theta/2}(\sqrt{2}P_*)^\theta = 2^{\theta/2}P_* \le \sqrt{2}P_*$. At $q = 0.9$, $t_0 = 42$, 63, and 84 for $n = 10^2$, 10^3, and 10^4 [dots in Fig. S1].

Proposition S6 (Memory bound)**.** *Let δ be the rate of Proposition S4.*

(a) Two reservoirs started in arbitrary states ρ and σ and driven by the same inputs obey $\|\Phi_{u_L}\cdots\Phi_{u_1}(\rho-\sigma)\|_1 \le 2(1-\delta)^{L-t_0}$ for $L \ge t_0$.

(b) Every input history h that extends into the infinite past determines a unique echo state $\rho(h)$, whose purity is at most P_*.

(c) If two input histories h and h' agree on their last L inputs, then $\frac{1}{2}\|\rho(h)-\rho(h')\|_1 \le (1-\delta)^L \le e^{-\delta L}$.

Proof. (a) The difference $X_t = d(\rho_t - \sigma_t)$ is traceless and evolves by the same maps as the states, $X_t = \mathcal{N}_q\mathcal{E}_{u_t}X_{t-1}$. After t_0 time steps, both states have purity at most $\sqrt{2}P_*$, so $|X_{t_0}|_1 \le 2$ and $|X_{t_0}|_2 \le 2\sqrt{2}P_*$, that is, $V(X_{t_0}) \le 2$. Proposition S4 then gives $|X_L|_1 \le V(X_L) \le 2(1-\delta)^{L-t_0}$. (b) Start the reservoir at earlier and earlier times in arbitrary states. By (a), the states it reaches at a given time converge, and the limit does not depend on the initialization; this limit is the echo state. By Eq. (S13), its purity is at most P_*. (c) Both echo states arise from earlier echo states, which belong to the histories before the last L inputs, through the same L time steps. The earlier echo states have purity at most P_*, so their difference has $V \le 2$ without any washout, and Proposition S4 applied L times gives the claim. □

Proof of Theorem 1(i). By Proposition S6(c) with $b = k$, two echo states whose input histories agree on their last L inputs are within trace distance ε as soon as $e^{-\delta L} \le \varepsilon$. Hence $W(\varepsilon) \le \lceil \delta^{-1}\ln(1/\varepsilon)\rceil$, and Corollary S5 gives Eq. (4) for $k \ge 4$, and the same bound with $k+1$ in place of k for $k \ge 1$. The argument uses the reservoir unitary, the input encoding, and the number of qubits only through the budget, so the bound holds for every reservoir. □

What the memory depth bounds. For two input histories that agree on their last L inputs, a linear readout $\mathrm{Tr}(F\rho)$ differs by at most $\|F\|\,\|\rho(h)-\rho(h')\|_1 \le 2\|F\|e^{-\delta L}$, and the best measurement tells the two histories apart with success probability at most $\frac{1}{2}(1+e^{-\delta L})$. Inputs older than $W(\varepsilon)$ thus change a readout of norm $\|F\|$ by at most $2\varepsilon\|F\|$. Large readout weights extend the reach only logarithmically: a readout with $\|F\| \le B$ sees inputs at most $\lceil \delta^{-1}\ln(2B/\varepsilon)\rceil$ time steps back at accuracy ε, while resolving a change of size ε in $\mathrm{Tr}(F\rho)$ from single-shot outcomes, whose variance is up to B^2, takes of order $(B/\varepsilon)^2$ repetitions.

Two further consequences follow directly. By (a), a reservoir started in an arbitrary state is within trace distance ε of the echo state at all times $t \ge t_0 + \lceil \delta^{-1}\ln(1/\varepsilon)\rceil$. And a reservoir with purely unitary input encoding has a budget of zero, so every time step maps the completely mixed state to itself; by (b), the completely mixed state is then the echo state for every input history, consistent with the observation that contractive reservoirs with input-independent fixed points carry no information about their inputs [16].

G. How tight are the bounds?

Figure S2(a) compares the two exponents. The exponent of the construction is the exact rate of the majority vote [Fig. S2(b)], so a sharper analysis of the construction cannot close the gap; the upper bound is conservative, for two reasons. Lemma S3 is a worst case over all differences with a given concentration, and its power $a = 2q/(1-q)$ grows quickly as the noise weakens. Moreover, V treats the 2-norm as if all of a reservoir's purity could be used to protect a single input. At $q = 0.5$, the upper bound grows by a factor of $2^{1.67}$ per input qubit, $\delta^{-1} = 2^{1.67k+6}$, and the construction by a factor of $e^{0.144}$; at $q = 0.9$, the factors are $2^{85.7}$ and $e^{0.83}$ [Fig. S2(c)]. The scaling with k is nevertheless the same, exponential on both sides, and neither bound depends on n.

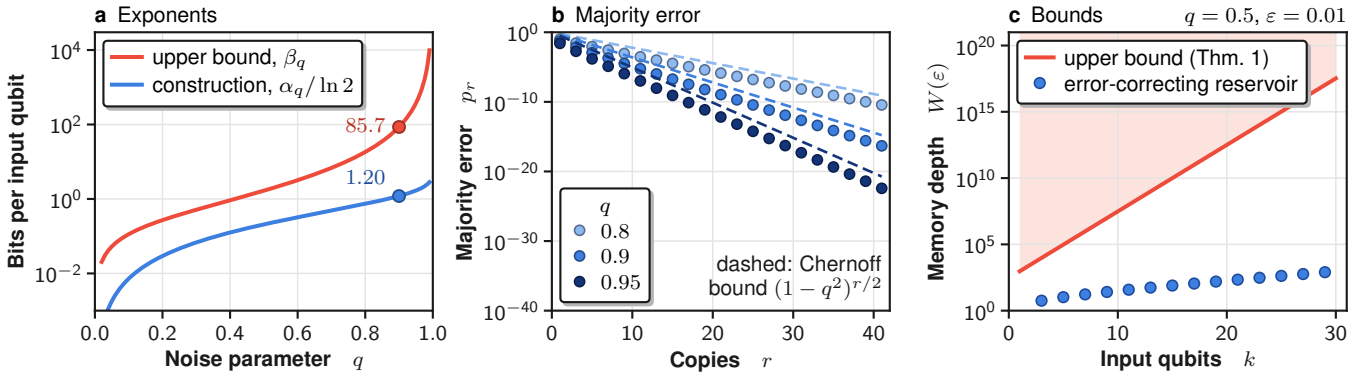


FIG. S2. **How tight are the bounds?** (a) Exponents per input qubit, in bits: β_q of the upper bound and $\alpha_q/\ln 2$ of the error-correcting reservoir; dots mark $q = 0.9$. (b) Exact majority error p_r (dots) and the Chernoff bound $(1-q^2)^{r/2}$ (dashed): the bound has the exact exponent. (c) Upper bound of Theorem 1 with the rate (S12) (line; the shaded region is excluded for every reservoir) and the memory depth of the error-correcting reservoir (dots; single-qubit readout) at $q = 0.5$. Both grow exponentially with k, at different rates.

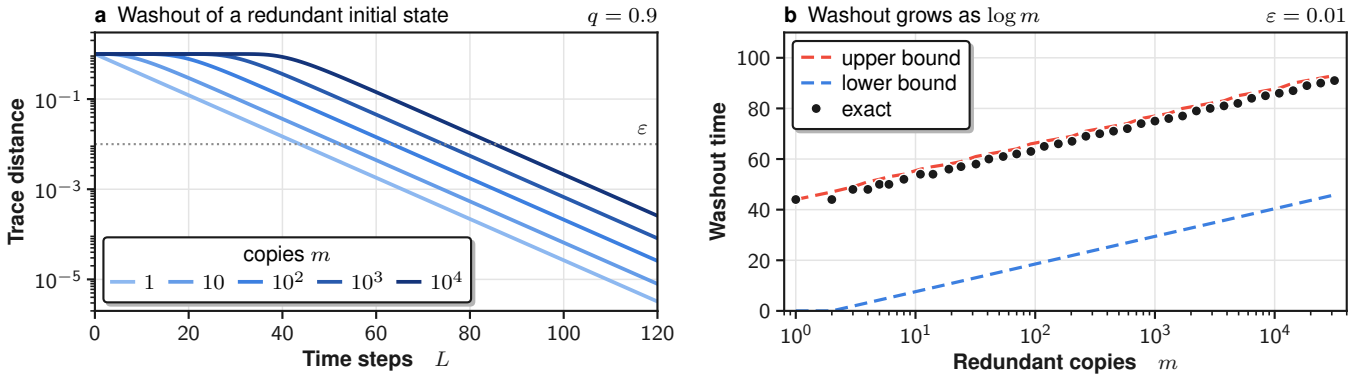


FIG. S3. **The washout grows with the reservoir size.** A reservoir with $U = I$ whose $m = n - k$ non-input qubits start in $|0\rangle^{\otimes m}$ or $|1\rangle^{\otimes m}$, at $q = 0.9$. (a) Exact trace distance between the two runs: more copies keep the initial state distinguishable for longer, but every curve eventually decays at the single-qubit rate q. (b) Exact washout time to $\varepsilon = 0.01$ (dots) and the two bounds derived in the text; all three grow as $\log m$.

IV. THE WASHOUT GROWS WITH THE RESERVOIR SIZE

A large reservoir can hold much information about its initial state, which new inputs cannot use. Take $U = I$, and let the $m = n - k$ qubits that are never reinitialized start either in $|0\rangle^{\otimes m}$ or in $|1\rangle^{\otimes m}$. After L time steps, every such qubit still shows its initial value with probability $(1 + q^L)/2$, independently of the others. The trace distance between the two runs therefore equals the total-variation distance between two binomial distributions, which we evaluate exactly [Fig. S3(a)]. Two simple bounds bracket it.

Lower bound. The sum of the Z outcomes has mean $\pm mq^L$ and variance at most m. By Chebyshev's inequality, deciding by majority errs with probability at most $1/(mq^{2L})$ under either initialization, so the trace distance is at least $1 - 2/(mq^{2L})$. It stays above ε as long as $mq^{2L} > 2/(1-\varepsilon)$, that is, for $L < \ln[m(1-\varepsilon)/2]/(2|\ln q|)$. A washout that grows as $\log m$ is therefore unavoidable.

Upper bound. The two product states have fidelity $(1 - q^{2L})^{m/2}$, so their trace distance is at most $\sqrt{m}\, q^L$. It falls below ε after $\lceil \ln(\sqrt{m}/\varepsilon)/|\ln q| \rceil$ time steps: copies that are never repaired become indistinguishable after $O_q(\log m)$ time steps.

Both bounds and the exact washout time grow as $\log m$ [Fig. S3(b)]. The same estimate shows that r unrepaired copies of an input keep it readable for only $O_q(\log r)$ time steps. The memory of new inputs, in contrast, is bounded independently of n.

V. THE ERROR-CORRECTING RESERVOIR

This section derives the exact response of the reservoir of the main text and the lower bound of Theorem 1, including the origin of the exponent α_q.

A. The construction

The reservoir has $n = 2r+2$ qubits: an input register of $k = r+2$ qubits, which consists of two input qubits (c, b) and r ancillas, and r memory qubits. A blank input, which asks the reservoir to keep its memory, injects $|c, b\rangle = |0, 0\rangle$; a store input for the bit b injects $|1, b\rangle$; the ancillas are always reinitialized in $|0\rangle^{\otimes r}$. The reservoir unitary is the permutation of computational basis states

$$U|z, c, b, x\rangle = |x,\, c,\, b,\, z \oplus F_{c,b}(x)\mathbf{1}\rangle, \tag{S14}$$

where z and x are the contents of the ancillas and of the memory qubits, $\mathbf{1} = (1, \ldots, 1)$, $F_{0,b}(x) = \mathrm{maj}(x)$, and $F_{1,b}(x) = b$. In words, U writes either the majority of the stored copies (blank input) or the new bit (store input) into the fresh ancillas, and it swaps ancillas and memory qubits. The map is invertible, because its output reveals x, c, and b, hence $F_{c,b}(x)$, and hence z; a permutation of basis states is a unitary. The construction is classical reversible logic run on qubits, and it uses no measurements and no feedback: majority-vote error correction of a repetition code [36], carried out unitarily with fresh ancillas, which is measurement-free error correction [26] in its simplest form.

B. Exact response

Under depolarizing noise, a qubit in $|0\rangle$ or $|1\rangle$ flips with probability $p = (1-q)/2$ and is otherwise left alone. Diagonal states thus stay diagonal, and every qubit behaves as a classical bit. Let $\hat{b}_j$ denote the bit that the unitary writes into the ancillas at time step j. During a blank time step, $\hat{b}_j = \mathrm{maj}(x)$, where x are the copies written at time step $j-1$, each of which has flipped with probability p during one round of noise, independently of the others. Hence $\Pr[\hat{b}_j \neq \hat{b}_{j-1}] = p_r$ for either value of $\hat{b}_{j-1}$, and

$$\mathbb{E}\big[(-1)^{\hat{b}_j}\big] = \lambda_r\, \mathbb{E}\big[(-1)^{\hat{b}_{j-1}}\big], \qquad \lambda_r = 1 - 2p_r. \tag{S15}$$

A stored bit sets $\hat{b} = b$ exactly, after which the bias decays as λ_r^t. Each memory qubit equals the current bit $\hat{b}$ up to one round of noise, which contributes the factor q in Eq. (7).

C. The majority error and the exponent α_q

The majority of r copies, for odd r, is wrong if more than half of them have flipped. Write $\xi_i = +1$ if copy i is intact and $\xi_i = -1$ if it has flipped, so that the ξ_i are independent with mean q. For every $\eta > 0$, Markov's inequality applied to $e^{-\eta \sum_i \xi_i}$ gives the Chernoff bound

$$p_r = \Pr\Big[\sum_i \xi_i < 0\Big] \le \Big[\tfrac{1+q}{2} e^{-\eta} + \tfrac{1-q}{2} e^{\eta}\Big]^r. \tag{S16}$$

The choice $e^{2\eta} = (1+q)/(1-q)$ minimizes the bracket, which then equals $\sqrt{1-q^2}$. Hence

$$p_r \le (1-q^2)^{r/2} = e^{-\alpha_q r}, \qquad \alpha_q = \tfrac{1}{2} \ln \frac{1}{1-q^2}, \tag{S17}$$

the bound quoted in the main text. The exponent α_q is the rate at which the majority vote becomes reliable as copies are added. It is exact, since $p_{r+2}/p_r \to 1 - q^2$ for large r; the Chernoff bound overestimates p_r only by a factor polynomial in r [Fig. S2(b)].

By Bernoulli's inequality, $\lambda_r^t \ge 1 - 2tp_r \ge \frac{1}{2}$ for $t \le e^{\alpha_q r}/4$, so the half-life of a stored bit is at least $\lfloor e^{\alpha_q r}/4 \rfloor$. Figure S4 shows the exact decay and half-lives at $q = 0.9$, and Table S1 lists exact half-lives for several noise strengths. From $k = 3$ to $k = 17$, the half-lives at $q = 0.9$ are 6, 47, 298, 1790, 10 431, 59 739, 337 940, and 1 894 286 time steps: seventeen input qubits per time step keep a stored bit readable for almost two million time steps, and every two additional input qubits multiply this time by more than five.

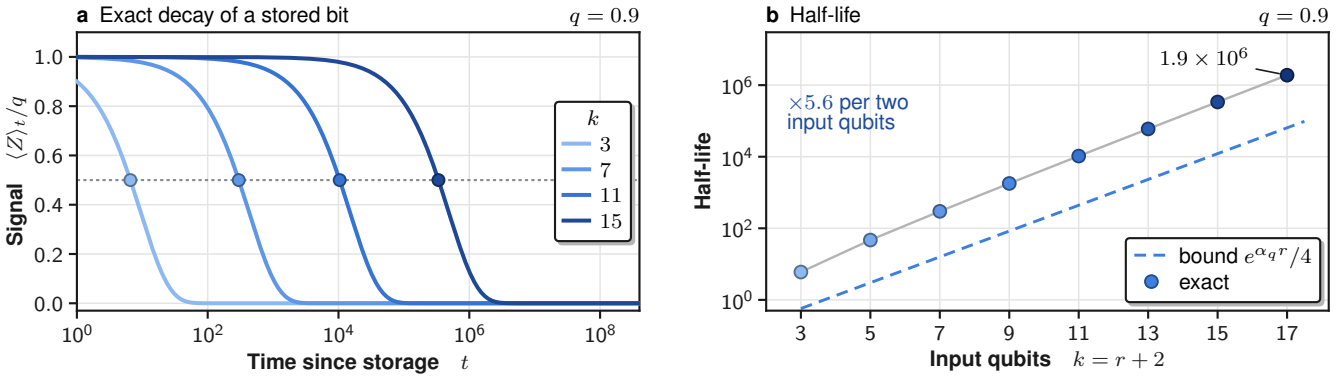


FIG. S4. **Exact response of the error-correcting reservoir** at $q = 0.9$. (a) Decay of a stored bit, $\langle Z\rangle_t/q = \lambda_r^t$; dots mark the half-lives. (b) Exact half-life versus the number of input qubits (dots) and the bound $e^{\alpha_q r}/4$ (dashed). Every two input qubits multiply the half-life by a nearly constant factor.

TABLE S1. Exact half-lives, in time steps, of a bit stored in the error-correcting reservoir: the largest T with $\lambda_r^T \geq 1/2$, for $k = r + 2$ input qubits.

q	$k = 3$	$k = 5$	$k = 9$	$k = 13$	$k = 17$
0.8	3	12	126	1171	10 306
0.9	6	47	1790	59 739	1.9×10^6
0.95	13	187	26 931	3.4×10^6	4.1×10^8
0.99	68	4636	1.6×10^7	4.9×10^{10}	1.4×10^{14}

D. Proof of Theorem 1(ii)

Two input histories that differ only in the bit stored L time steps ago, followed by blanks, agree on their last L inputs. Because a stored bit overwrites the memory qubits regardless of their previous content, both histories lead to valid echo states, and the memory is not a remnant of the initialization. By Eq. (7), the two echo states give expectation values of Z on a memory qubit that differ by $2q\lambda_r^L$. Since $|\operatorname{Tr}[O(\rho - \rho')]| \leq \|O\| \, \|\rho - \rho'\|_1$, their trace distance is at least $q\lambda_r^L$, which exceeds ε as long as $L < \ln(q/\varepsilon)/(-\ln \lambda_r)$. Hence

$$W(\varepsilon) \geq \frac{\ln(q/\varepsilon)}{-\ln \lambda_r} \geq \frac{1}{4} e^{\alpha_q r} \ln \frac{q}{\varepsilon}. \tag{S18}$$

The first inequality requires $\varepsilon < q$, because the readout qubit's own last round of noise contributes the factor q. For the second, we use $-\ln(1-x) \leq 2x$ for $0 \leq x \leq \frac{1}{2}$ with $x = 2p_r$; this requires $p_r \leq \frac{1}{4}$, and then $-\ln \lambda_r \leq 4p_r \leq 4e^{-\alpha_q r}$ by Eq. (S17).

For a general number k of input qubits, choose r as the largest odd number with $r \leq k-2$, and leave the remaining input qubits unused. Then $r \geq k - 3$ and $n = k + r \leq 2k$. If $e^{\alpha_q (k-3)} \geq 4$, then $p_r \leq e^{-\alpha_q r} \leq \frac{1}{4}$, and Eq. (S18) gives Eq. (5). At $q = 0.9$, the condition holds for $k \geq 5$. □

E. Further properties

The single-qubit readout is nearly optimal. The lower bound uses the readout of a single memory qubit. The full reservoir state carries slightly more information: for $k = 5$ and $q = 0.9$, the exact trace distance between the two echo states [Fig. S6(a)] gives a memory depth of 316 time steps at $\varepsilon = 0.01$, compared with 309 for the single-qubit readout of Eq. (S18).

Without repair. The same reservoir without the majority vote writes the bit into its r memory qubits and then leaves them alone. After t further time steps, every copy has the Bloch vector $\pm q^{t+1}\hat{z}$, and the trace distance between the two echo states is the binomial sum $\frac{1}{2}\sum_{j=0}^{r} \binom{r}{j} \left| a^j b^{r-j} - b^j a^{r-j} \right|$ with $a, b = (1 \pm q^{t+1})/2$, which no readout can beat. At $q = 0.9$ and $\varepsilon = 0.01$, the memory depth of the copies grows only from 43 time steps for $r = 1$ to 53 for

TABLE S2. Entropy budget of common ingredients of quantum reservoirs: how many bits of purity a process can inject per time step. The memory depth is at most exponential in the total budget, and processes with zero budget cannot raise the memory limit. Processes outside the model are not covered by the bound.

Process	Budget per time step	Consequence for the memory
Reinitialization of k input qubits	k bits	memory depth $e^{\Theta(k)}$ under depolarizing noise (Theorem 1)
Measurement followed by reset of m qubits	m bits	counts like m reinitialized input qubits [11, 12]
Unitary input encoding (input-dependent gates)	0	without another entropy sink, the echo state is the completely mixed state
Projective (or other Lüders) measurement of reservoir qubits whose outcomes only enter the readout	0	the averaged measurement channel is unital and cannot raise the memory limit [33]; general instruments count with their own budget, and indirect measurements through fresh meter qubits count like reinitialized qubits [17]
Unital noise (Pauli errors, dephasing, coherent errors, crosstalk)	0	leaves the memory limit unchanged
Amplitude damping with probability $0 < \gamma < 1$ on all n qubits	$n \log_2(1+\gamma)$ bits	memory depth can grow exponentially with n, already for $k = 2$ [13, 21, 23]
Feedback of measurement outcomes; non-Markovian noise; noise without a depolarizing component	outside the model	memory can be unbounded (Sec. VI E) [18]

$r = 13$, whereas the error-correcting reservoir with $r = 13$ copies ($k = 15$) remembers 2.2×10^6 time steps [Fig. 4(b)]. The fresh qubits extend the memory only because the reservoir uses them for repair.

Retention cost. Suppose that a reservoir must keep the expectation values of an observable O with $\|O\| \leq 1$ for a stored 0 and a stored 1 at least q apart for T time steps. Proposition S6(c) bounds this gap by $2e^{-\delta T}$, so $T \leq \delta^{-1} \ln(2/q)$, and with $\delta^{-1} \leq 2^{k/(1-q)^2}$, $k = \Omega_q(\log T)$. Conversely, Eq. (7) keeps the gap for all $t \leq T$ if $T \leq e^{\alpha_q r}/4$. An odd $r \geq \ln(4T)/\alpha_q$ and $k = r + 2$ therefore suffice, so $k = O_q(\log T)$. Retaining an input for T time steps thus requires, and can be achieved with, $\Theta_q(\log T)$ reinitialized qubits per time step.

Anisotropic noise. For independent unital noise $\mathcal{M}$ on every qubit, the bit is best stored along the Bloch axis that $\mathcal{M}$ preserves best. Let $s_{\mathcal{M}}$ be the largest singular value of the matrix by which $\mathcal{M}$ acts on Bloch vectors, with unit vectors $\boldsymbol{v}$ and $\boldsymbol{v}'$ such that $\mathcal{M}$ maps the Bloch vector $\boldsymbol{v}$ to $s_{\mathcal{M}}\boldsymbol{v}'$. Conjugating the memory block of Eq. (S14) by single-qubit rotations from the $\boldsymbol{v}'$ basis and into the $\boldsymbol{v}$ basis makes every memory qubit an independent copy of the stored bit, flipped with probability $(1 - s_{\mathcal{M}})/2$. The lower bound then holds with q replaced by $s_{\mathcal{M}}$.

Non-Clifford dynamics. The majority vote is a non-Clifford operation, as long-term memory requires [37].

Memory is not computation. A long memory does not imply computational power or classical hardness: a single number, the bias λ_r^t, describes the entire response of this reservoir.

VI. ENTROPY BUDGET BEYOND DEPOLARIZING NOISE

This section extends Theorem 1 to more general noise, using the budget of Sec. II, and shows where the extension ends. Table S2 summarizes the rule.

A. Unital noise and measurements

Let the noise at time step t be $\mathcal{R}_t \circ \mathcal{N}_q \circ \mathcal{S}_t$, where $\mathcal{R}_t$ and $\mathcal{S}_t$ are unital channels. They can be correlated across qubits, coherent, input dependent, or drifting in time. Two observations reduce this case to the one already proven. First, $\mathcal{S}_t$ can be absorbed into the injection step: $\mathcal{S}_t \circ \mathcal{E}_u$ is a channel, and because $\mathcal{S}_t$ is positive and unital, $\mathcal{S}_t\mathcal{E}_u(I) \leq 2^k \mathcal{S}_t(I) = 2^k I$, so it has the same budget of k bits. Second, $\mathcal{R}_t$ cannot increase V, because a unital channel contracts both norms (Lemma S1 with $b = 0$). Hence $V(\mathcal{R}_t\mathcal{N}_q\mathcal{S}_t\mathcal{E}_u X) \leq V(\mathcal{N}_q\mathcal{S}_t\mathcal{E}_u X) \leq (1-\delta)V(X)$, and the same two facts give the purity recursion. Two input histories that agree on their last L inputs share every map during those L time steps, so Theorem 1 holds verbatim; its lower bound needs only $\mathcal{R}_t = \mathcal{S}_t = \mathrm{id}$.

Measurements are covered in the same way. A measurement whose post-measurement states follow the Lüders rule, with Kraus operators $\sqrt{E_i}$, is unital when its outcomes are discarded, because $\sum_i \sqrt{E_i}\sqrt{E_i} = \sum_i E_i = I$; this includes every projective measurement. Such measurements of the reservoir qubits whose outcomes only enter the readout therefore cannot raise the memory limit. A general measurement instrument with Kraus operators A_i, $\sum_i A_i^\dagger A_i = I$, is unital only if also $\sum_i A_i A_i^\dagger = I$; the Kraus operators of amplitude damping are a counterexample. A general instrument therefore counts with its budget $b = \log_2 \|\sum_i A_i A_i^\dagger\|$, like any other channel. An indirect measurement through a meter qubit that is freshly prepared at every time step is of this kind: the meter is a reinitialized qubit and adds one bit of budget per time step, consistent with the role of measurement back-action as a source of dissipation [17].

B. Non-unital noise and resets

Now let $\mathcal{R}_t$ and $\mathcal{S}_t$ be arbitrary channels with budgets of b_R and b_S bits, with $\mathcal{R}_t$ independent of the inputs. Regroup the time steps as $\rho_t = \mathcal{R}_t(\tilde{\rho}_t)$, with

$$\tilde{\rho}_t = \Psi_t(\tilde{\rho}_{t-1}), \qquad \Psi_t = \mathcal{N}_q \circ \mathcal{S}_t \circ \mathcal{E}_{u_t} \circ \mathcal{R}_{t-1}, \tag{S19}$$

and $\mathcal{R}_0 = \mathrm{id}$. By positivity, $\mathcal{S}_t\mathcal{E}_{u_t}\mathcal{R}_{t-1}(I) \le 2^{k+b_R+b_S} I$, so Sec. III applies to $\tilde{\rho}_t$ with the budget $b = k + b_R + b_S$. Two input histories that agree on their last L inputs share every Ψ_t during those L time steps. The final channel $\mathcal{R}_t$ cannot increase trace distances, so the memory depth of ρ_t is at most that of $\tilde{\rho}_t$. Apart from the washout time, the reservoir size thus enters only through $b_R + b_S$. The bounds hold for every realization of random channels $\mathcal{R}_t$ and $\mathcal{S}_t$, and hence also on average.

Two budgets are used in the main text. A reset channel $\rho \mapsto \mathrm{Tr}(\rho)|0\rangle\langle 0|$ maps I to $2|0\rangle\langle 0| \le 2I$, so a reset of m qubits has a budget of m bits. Amplitude damping $\mathcal{AD}_\gamma$ with probability γ maps I to $\mathrm{diag}(1+\gamma, 1-\gamma) \le (1+\gamma)I$, so a layer of amplitude damping on all n qubits has a budget of $n \log_2(1+\gamma)$ bits.

C. Every unital qubit channel with full-rank Choi matrix contains depolarizing noise

Let $\mathcal{M}$ be a unital qubit channel, and let $\Omega_\mathcal{M} = \sum_{i,j=0}^{1} \mathcal{M}(|i\rangle\langle j|) \otimes |i\rangle\langle j|$ be its Choi matrix, a positive operator with $\mathrm{Tr}\,\Omega_\mathcal{M} = 2$. Complete depolarization, $\mathcal{D}_0(X) = \mathrm{Tr}(X) I/2$, has $\Omega_{\mathcal{D}_0} = I/2$. Suppose that $\Omega_\mathcal{M}$ is positive definite and that $\mathcal{M} \neq \mathcal{D}_0$. Then the smallest eigenvalue of $\Omega_\mathcal{M}$ lies in $(0, 1/2)$, and $q_\mathcal{M} = 1 - 2\lambda_{\min}(\Omega_\mathcal{M}) \in (0, 1)$. The map $\mathcal{T} = [\mathcal{M} - (1 - q_\mathcal{M})\mathcal{D}_0]/q_\mathcal{M}$ is trace preserving and unital, and its Choi matrix $[\Omega_\mathcal{M} - \lambda_{\min}(\Omega_\mathcal{M}) I]/q_\mathcal{M}$ is positive semidefinite, so $\mathcal{T}$ is a channel. Because $\mathcal{T} \circ \mathcal{D}_q = q\mathcal{T} + (1-q)\mathcal{D}_0$ for every unital channel $\mathcal{T}$, we obtain

$$\mathcal{M} = \mathcal{T} \circ \mathcal{D}_{q_\mathcal{M}}. \tag{S20}$$

Every such channel is thus depolarizing noise of strength $q_\mathcal{M}$ followed by a unital channel, which is the form covered above. For a Pauli channel with probabilities p_I, p_X, p_Y, p_Z, the Choi eigenvalues are $2p_I, \ldots, 2p_Z$, so $1 - q_\mathcal{M} = 4\min\{p_I, p_X, p_Y, p_Z\}$. For example, Pauli noise with $p_X = p_Y = 0.01$ and $p_Z = 0.08$ has $q_\mathcal{M} = 0.96$.

D. Cooling noise: memory exponential in n

Amplitude damping can make the memory grow with the reservoir size, already with two input qubits. The construction uses idle qubits as a refrigerator [23]. Let every time step end with the noise $(\mathcal{D}_q \circ \mathcal{AD}_\gamma)^{\otimes n}$, with $0 < \gamma < 1$.

Idle qubits cool. On diagonal states, one time step of noise acts on every qubit as a bit channel: a qubit that is 1 with probability x becomes 1 with probability $p + \kappa x$, where $p = (1-q)/2$ and $\kappa = q(1-\gamma)$. Here $0 < \kappa < 1$ is the fraction of a qubit's polarization that survives one time step. The condition $\gamma < 1$ is necessary: for $\gamma = 1$, the noise resets every qubit to $|0\rangle$ before depolarizing it, the state after every time step is $[(I + qZ)/2]^{\otimes n}$ whatever the inputs, and the memory depth is zero. A qubit left alone approaches the stationary probability $p_\infty = p/(1-\kappa)$ of being 1, which is below $1/2$ for every $\gamma > 0$: idle qubits cool. After s time steps, starting from either value, a qubit is 1 with probability at most $\mu_s = p_\infty + \kappa^s(1 - p_\infty)$, and we choose the cooling time s such that $\mu_s < 1/2$.

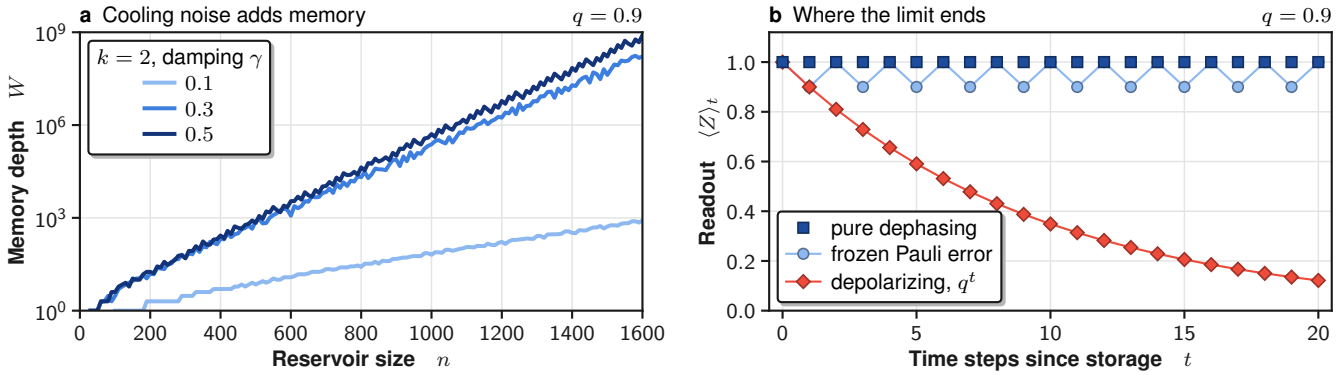


FIG. S5. **The entropy budget at work.** (a) Cooling noise adds memory. Lower bound $\lfloor 1/(4p_{\rm f}) \rfloor$ on the memory depth of the cooling reservoir of Sec. VI D, with only $k = 2$ input qubits, under the noise $(\mathcal{D}_q \circ \mathcal{AD}_\gamma)^{\otimes n}$; $p_{\rm f}$ is the worst-case probability of losing the stored bit per time step, from exact binomial tails. (b) Where the limit ends. Readout of a stored bit under pure dephasing (no decay), under a Pauli error that is drawn once and then repeated at every time step (perfect revivals, although one time step is depolarizing on average), and under Markovian depolarizing noise (q^t).

The reservoir. The reservoir has the two input qubits (c, b) of Sec. V and $s + 1$ blocks $x_0, \ldots, x_s$ of $r = \lfloor (n - 2)/(s + 1) \rfloor$ qubits each. Block x_0 holds the stored bit, and the other blocks form a queue in which qubits cool. The reservoir unitary is the basis permutation

$$U|c, b\rangle|x_0, x_1, \ldots, x_s\rangle = |c, b\rangle|x_s \oplus F_{c,b}(x_0)\mathbf{1},\, x_0, \ldots, x_{s-1}\rangle. \tag{S21}$$

For a store input, $F_{1,b}(x) = b$. For a blank input, $F_{0,b}(x) = 1$ if the fraction of ones in x is at least the threshold $\chi = p + \kappa/2$, and $F_{0,b}(x) = 0$ otherwise. In words, U decodes the stored bit by a threshold vote, writes it into the coolest block x_s, and moves the other blocks one position along the queue.

Error per time step. The reused block has cooled for s time steps, so each of its bits differs from the newly written bit F with probability at most μ_s, independently given the past. After the noise, each bit of the new memory block is therefore 1 with probability at least $p + \kappa(1 - \mu_s)$ if $F = 1$, and at most $p + \kappa\mu_s$ if $F = 0$. The threshold χ lies between these two values, at a distance of at least $\kappa(1/2 - \mu_s)$ from both. The next vote therefore errs with probability at most

$$p_{\rm f} = \max\big\{\Pr[\mathrm{Bin}(r, p + \kappa\mu_s) \geq \chi r],\ \Pr[\mathrm{Bin}(r, p + \kappa(1 - \mu_s)) < \chi r]\big\}, \tag{S22}$$

and Hoeffding's inequality [34] gives $p_{\rm f} \leq e^{-\nu r}$ with $\nu = 2\kappa^2(1/2 - \mu_s)^2$. The failure probability thus falls exponentially with the block size.

Memory. Consider two input histories that differ only in the bit stored L time steps ago, followed by blanks. The stored bit survives the $L + 1$ subsequent votes with probability at least $1 - (L + 1)p_{\rm f}$, so the threshold observable of block x_0, $+1$ if the fraction of ones is at least χ and -1 otherwise, differs between the two echo states by at least $2 - 4(L + 1)p_{\rm f}$. For $(L + 1)p_{\rm f} \leq 1/4$, this difference is at least 1, and the trace distance is at least $1/2$. Hence

$$W(\varepsilon) \geq \left\lfloor \frac{1}{4p_{\rm f}} \right\rfloor \geq \left\lfloor \frac{e^{\nu r}}{4} \right\rfloor \qquad \text{for all } \varepsilon < \frac{1}{2}. \tag{S23}$$

With only two input qubits, the memory depth is exponential in $r \approx n/(s + 1)$, and hence in n [Fig. S5(a)].

This is the classical counterpart of the quantum refrigerator of Ref. [23]; for fault-tolerant quantum computation under amplitude damping, the achievable length is known to be at most exponential in n [27].

Numbers. At $q = 0.9$ and $\gamma = 0.3$, we have $p = 0.05$, $\kappa = 0.63$, and $p_\infty \approx 0.135$. The exact binomial tails, optimized over the cooling time s, push $p_{\rm f}$ below 10^{-6} at $n = 1010$ qubits, with $s = 7$ and $r = 126$. The cooling route is expensive: without damping, the error-correcting reservoir reaches a failure probability of 1.8×10^{-7} per time step with 32 qubits, 17 of them input qubits. On hardware dominated by energy relaxation, the reservoir size does matter for memory, through cooling, consistent with the memory gains attributed to amplitude damping in gate-based reservoirs [13] and with collective relaxation as a scalable fading-memory mechanism [21].

E. Where the limit ends: two counterexamples

Two assumptions of Theorem 1 cannot be dropped: that the noise has a depolarizing component, and that it is Markovian. Consider a single memory qubit next to the two input qubits (c, b). The reservoir unitary swaps b into the memory qubit when $c = 1$ and acts trivially when $c = 0$; the old memory value leaves with the input register at the next input injection.

Pure dephasing. Under the dephasing channel $\mathcal{Z}_\phi(X) = (1-\phi)X + \phi\, ZXZ$, computational-basis states are fixed points, so a stored bit stays perfectly readable forever.

Frozen Pauli errors. Let each qubit carry a Pauli operator $S \in \{I, X, Y, Z\}$, drawn once with probabilities $\frac{1+3q}{4}, \frac{1-q}{4}, \frac{1-q}{4}, \frac{1-q}{4}$ and then applied at every time step. Averaged over the draw, one time step is exactly $\mathcal{D}_q^{\otimes n}$, because

$$\mathbb{E}[S\rho S] = \tfrac{1+3q}{4}\rho + \tfrac{1-q}{4}(X\rho X + Y\rho Y + Z\rho Z) = \mathcal{D}_q(\rho). \tag{S24}$$

But $S^2 = I$, so after an even number of time steps the memory qubit is back in $|b\rangle\langle b|$, and the stored bit is read perfectly at every second time step. Averaged over the draw, the readout alternates between 1 and q [Fig. S5(b)]. Here the environment remembers the noise, which is why the theorem requires Markovian noise; non-Markovian dynamics can likewise extend the memory of quantum reservoirs in practice [18].

In both cases, the reservoir has no unique echo state, and no finite memory depth exists for input histories that contain a stored bit. Finally, a reservoir whose measurement outcomes are stored and fed back is coupled to a classical memory outside the model, which is an entropy sink of its own.

VII. NUMERICAL SIMULATIONS

A. Reservoirs

We simulate the density matrix of Eq. (2) exactly for $n \le 9$ qubits. The first k qubits form the input register. Each input qubit is reinitialized in $|\psi(u_t)\rangle = \sqrt{1-u_t}\,|0\rangle + \sqrt{u_t}\,|1\rangle$ [6], with independent inputs u_t drawn uniformly from $[0, 1]$. We consider two reservoir unitaries. The *Ising reservoir* evolves for a time Δt under the transverse-field Ising Hamiltonian

$$H = \sum_{i<j} J_{ij} X_i X_j + g \sum_i Z_i, \tag{S25}$$

with couplings J_{ij} drawn uniformly from $[-J/2, J/2]$, transverse field $g = J$, and $J\Delta t = 10$, as in the standard model of QRC [6, 7, 32]. The *Haar reservoir* uses a unitary drawn from the Haar measure on $U(2^n)$, the extreme of information scrambling. After every time step, each qubit undergoes depolarizing noise $\mathcal{D}_q$. Table S3 lists all parameters.

B. Memory depth

The memory depth of Eq. (3) is a supremum over all pairs of input histories that agree on their last L inputs. Because each time step is linear, the difference Δ_L of the two reservoir states evolves under the same map as the states themselves, and we follow it directly. We evaluate two kinds of history pairs for every reservoir. In the first, the histories differ in a single input: we prepare the echo state of a random input history and inject the two most different inputs, $u = 0$ and $u = 1$, into two copies of it. In the second, the histories differ in all earlier inputs: we prepare the echo states of the constant histories $u = 0$ and $u = 1$. Each pair is followed by five common continuations, two random input sequences and the constant inputs $u = 0$, $1/2$, and 1, and we record $D(L) = \frac{1}{2}\|\Delta_L\|_1$ after L common time steps until $D(L) < 3\times10^{-3}$, well below ε. Our estimate of a reservoir's memory depth is $1+\max\{L : D(L) > \varepsilon\}$, maximized over both pairs and all continuations; as a maximum over a subset of all pairs, it can only underestimate the supremum, and check (iv) below tests it against a much larger set.

Histories that differ in all earlier inputs set the memory depth of almost every reservoir: in the main runs [Figs. 4(c) and S7], they give a larger value than a single differing input in 94 of 180 cases, the same value in 81, and a smaller one in 5. Their differences accumulate in the slowest modes of the reservoir, which a single input excites only weakly; for the Ising reservoirs with one input qubit, they add about five time steps. The idle qubit, a single qubit that holds the input without repair, has memory depth $1 + \max\{L : q^{L+1} > \varepsilon\} = 43$ at $q = 0.9$ and $\varepsilon = 0.01$.

TABLE S3. Parameters of the simulations in Fig. 4 of the main text and Figs. S6 and S7.

Parameter	Value
Noise parameter q	0.9; also 0.95–1 [Fig. S6(b)]
Accuracy ε	0.01
Reservoir size n	3–9
Input qubits k	1–4
Ising couplings	$J_{ij} \in [-J/2, J/2]$, field $g = J$
Evolution time	$J\Delta t = 10$; also 0.3–30 [Fig. S6(c)]
Input encoding	$\sqrt{1-u}\,\|0\rangle + \sqrt{u}\,\|1\rangle$, $u \in [0,1]$; cold test: additional input qubits in $\|0\rangle$
Washout	until three initial states agree within 10^{-6} (60–2360 time steps)
History pairs	one differing input; all earlier inputs differing
Continuations	5 per history pair
Realizations	6 reservoirs per n and k, the same for all q, $J\Delta t$, and encodings
Stopping criterion	$D(L) < 3 \times 10^{-3}$

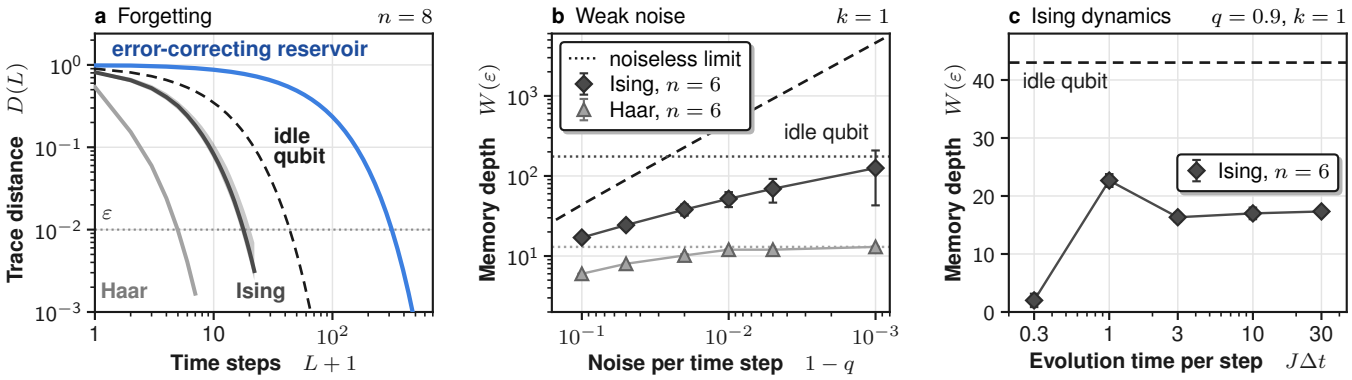


FIG. S6. **Simulations of standard reservoirs.** (a) Trace distance $D(L)$ between the echo states of two input histories that agree on their last L inputs, for reservoirs with $n = 8$ qubits: Haar and Ising reservoirs with one input qubit, for histories that differ in all earlier inputs (median and range over six reservoirs, first random continuation), an idle qubit (dashed), and the error-correcting reservoir with $k = 5$, for histories that differ in the stored bit (blue). The memory depth is where $D(L)$ falls below ε (dotted). (b) Memory depth of the same six reservoirs with $n = 6$ and $k = 1$ versus the noise per time step (mean and standard deviation), compared with an idle qubit; dotted lines mark the noiseless limit. (c) Memory depth of the six Ising reservoirs with $n = 6$ versus the evolution time per time step, with fixed couplings.

C. Validation

Four checks confirm the code and the estimator. (i) The evolution of differences agrees with the difference of evolved states to machine precision. (ii) For the error-correcting reservoir, the simulated single-qubit readout reproduces the exact response $q\lambda_r^L$ of Eq. (7) to machine precision for $r = 1$ and $r = 3$. (iii) Every echo state is prepared by a washout that continues until runs from the completely mixed state and from the product states $|0\cdots0\rangle$ and $|1\cdots1\rangle$, driven by the same inputs, agree within a trace distance of 10^{-6}, far below ε. This takes at most 130 time steps at $q = 0.9$ and up to 2360 time steps as $q \to 1$, including the noiseless limit, which the theorem does not cover; there, the three runs agree for every simulated reservoir. (iv) A much larger set of history pairs changes the memory depth by at most one time step. For 24 reservoirs (Ising and Haar, $n = 6$ and 8, $k = 1$ and 2), we compared the estimator with blocks of 2 to 30 differing inputs, all 32 sign patterns of six differing inputs up to exchanging the two histories, constant histories with the input pairs $(0, \frac{1}{2})$, $(\frac{1}{2}, 1)$, and $(\frac{1}{4}, \frac{3}{4})$, three pairs of independent random histories, and continuations that maximize the trace distance greedily at every time step.

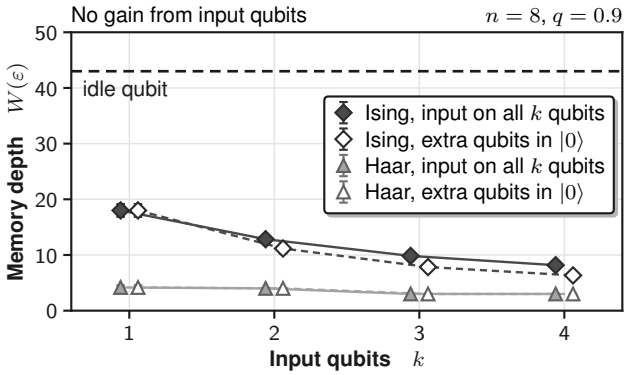


FIG. S7. **More input qubits do not help standard reservoirs.** Memory depth of the six Ising and the six Haar reservoirs with $n = 8$ versus the number of input qubits, when all input qubits carry the input (filled) and when only the first one does and the others are reinitialized in $|0\rangle$ (open; mean and standard deviation). For $k = 1$, both protocols coincide. Unitary ensembles are sampled separately for each k; both protocols share the ensemble at fixed k.

D. Results and robustness

The forgetting curves of Fig. S6(a) show the hierarchy of the main text directly: the Haar reservoir forgets within a few time steps, the Ising reservoir within about 18, the idle qubit within 43, and the error-correcting reservoir with the same $n = 8$ only after 316 time steps. The hierarchy persists for weaker noise and other evolution times. As the noise weakens, the memory depth of the same reservoirs grows far more slowly than that of an idle qubit, which grows as $1/(1-q)$, and it saturates in the noiseless limit [Fig. S6(b)]. There, the reinitialization of the input register is the only source of forgetting: every time step discards the part of the stored information that the dynamics has moved into the input register. Even a noiseless reservoir therefore forgets, within 75 to 539 time steps for our Ising reservoirs, 175 on average, and within 13 time steps for the Haar-random ones at $n = 6$. For the Ising reservoirs, the mean memory depth lies between 16 and 23 time steps for all evolution times from $J\Delta t = 1$ to 30 [Fig. S6(c)], far below the idle qubit. For very short evolution times, the input hardly leaves the input register, so the reservoir forgets it at the next injection.

Cold input qubits. In the runs above, every input qubit carries the input, so a larger k also overwrites more of the reservoir with new input. To separate the two effects, we repeated the runs at $n = 8$ for the same reservoirs with only the first input qubit carrying the input and the other $k - 1$ input qubits reinitialized in $|0\rangle$, which makes them pure entropy sinks. The memory depth still falls as k grows [Fig. S7], from 18 time steps for $k = 1$ to 11, 8, and 6 for $k = 2$, 3, and 4 (Ising, means over the six reservoirs), and from 4 to 4, 3, and 3 (Haar), compared with 13, 10, and 8 (Ising) when all input qubits carry the input. A cold qubit removes entropy, but a generic unitary moves stored information into it as readily as errors, and the next reinitialization erases both. Only dynamics that separates the two, as the error-correcting reservoir does, turns fresh qubits into memory.

VIII. IMPLICATIONS FOR RESERVOIR DESIGN, AND SCOPE

A. Design lessons

Return to the forecasting task of Fig. 2(a), $y_{t+1} = 0.7\,u_t + 0.3\,u_{t-100}$, and consider two reservoirs with the same noise, $q = 0.9$, and the same $k = 2$ input qubits per time step, a reservoir A with $n = 20$ qubits and a reservoir B with $n = 200$ qubits [Fig. 3(a)]. To forecast y_{t+1}, a reservoir must retain u_{t-100}, that is, it needs a memory depth of at least 100 time steps at an accuracy that resolves the second term. In fact, it must retain every input of the last 100 time steps at once, a stronger requirement for which the memory depth is necessary but not sufficient. The natural remedy for a reservoir that falls short is to enlarge it. Theorem 1 shows that this remedy has a limit that does not move with n: with two input qubits, the memory depth of every reservoir, A, B, or any larger one, is bounded by the same constant $W_{q,2}(\varepsilon)$, whereas every additional input qubit can multiply the attainable memory depth by a constant factor. Our constants are far too loose to decide whether 100 time steps lie below $W_{0.9,2}(\varepsilon)$; the theorem identifies the knob, not its setting. The simulations add a practical scale: standard Ising reservoirs remember 16 to 18 time

steps at this noise level, for every size from $n = 3$ to $n = 9$. Four lessons follow.

More qubits are not automatically better. A larger reservoir can generate richer features of its recent inputs [7], but at fixed k and q its best memory depth is bounded independently of n. Large random reservoirs can even suffer from an exponential concentration of their outputs, which makes different inputs hard to distinguish [14, 38], and they forget faster [Fig. 4(c)].

The input-injection architecture is decisive. The number of qubits reinitialized per time step sets the limit. More reinitialized qubits do not guarantee a longer memory, however, because reinitialization also erases whatever those qubits stored, even when the additional qubits are reinitialized in $|0\rangle$ [Fig. S7]. The limit is reached only if the reservoir dynamics moves the relevant information out of the input register before it is reinitialized and uses the fresh qubits to repair it, as the error-correcting reservoir does.

The memory requirement of the task matters. If a task depends on inputs further back than the memory depth $W(\varepsilon)$, these inputs change every trained readout by at most $2\varepsilon\|F\|$. Estimating how far back the task-relevant inputs lie is therefore a prerequisite for choosing k.

A long forecasting horizon is not a long memory. Predicting many steps ahead does not by itself require a deep memory; what matters is how far back the inputs lie that determine the target. In the task above, already a one-step forecast requires a memory depth of 100 time steps.

B. What the result does not say

Memory depth is necessary for a temporal task, but it is not sufficient. A reservoir must also transform its input history nonlinearly and make the relevant information visible to its readout [35], and the error-correcting reservoir, whose memory scaling is optimal, computes nothing beyond a delayed copy of its input. Theorem 1 therefore makes no statement about task performance, nonlinear processing capacity, sampling noise, or a quantum advantage. It bounds how far back any task can reach, not how well a reservoir solves tasks within that range. In particular, the theorem does not provide an optimal QRC architecture for time series. It does not tell how many input qubits a reservoir of 20 qubits should use, which couplings generate the most useful temporal features [32, 35], how much memory a given real-world time series requires, or when a quantum reservoir outperforms a classical one. The constants of the upper bound are conservative by many orders of magnitude, and our simulations cover sizes up to $n = 9$. The result is a fundamental design principle rather than a design recipe.

IX. CODE AND DATA

Project repository: `https://github.com/eybmits/fresh-qubits-memory`.

The source archive of this Supplemental Material contains the code and data behind every number, table, and figure. The folder `sim` holds the density-matrix simulator and the memory-depth estimator (`qrc_sim.py`), the simulation runs (`campaign.py`, `merge.py`), the validation of the code and of the estimator (`test_sim.py`, `validate_pairs.py`), the error-correcting reservoir (`rep_curve.py`), the exact constants, half-lives, washout times, and the cooling reservoir (`exact_numbers.py`), and all results as JSON files. The folder `figures` holds the sources of all figures, and the folder `checks` independent checks of the constants and of the elementary steps of the proof (`verify_constants.py`, `verify_simple_proof.py`). A README file lists the commands that rebuild every figure from these data. The code was written with the assistance of Claude Opus 5.5 (Anthropic) under the direction of Markus Baumann, who verified it with the checks of Sec. VII and against the exact results of Secs. IV and V.